

Fairness Aware Ranking & Recommendations in Networks

A Thesis

submitted to the designated
by the General Assembly
of the Department of Computer Science and Engineering
Examination Committee

by

Sotirios Tsioutsoulis

in partial fulfillment of the requirements for the degree of

MASTER OF SCIENCE IN DATA AND COMPUTER
SYSTEMS ENGINEERING

WITH SPECIALIZATION
IN DATA SCIENCE AND ENGINEERING

University of Ioannina

September 2020

Examining Committee:

- **Evaggelia Pitoura**, Professor, Department of Computer Science and Engineering, University of Ioannina (Advisor)
- **Panayiotis Tsaparas**, Associate Professor, Department of Computer Science and Engineering, University of Ioannina
- **Nikolaos Mamoulis**, Professor, Department of Computer Science and Engineering, University of Ioannina

DEDICATION

To my family.

ACKNOWLEDGEMENTS

I have to start by expressing my gratitude to my supervisor Prof. Evaggelia Pitoura for the guidance, the support and the valuable remarks and advices she provided me during the completion of this thesis. Furthermore, I would like to thank my co-supervisor Prof. Panayiotis Tsaparas for his significant contributions and the excellent collaboration. Last, I want to thank Prof. Nikolaos Mamoulis for being always willing to help me with his useful insights.

Special thanks to my friend and my fellow student Ilias Kleftakis for his remarks and his contribution in the implementation of the algorithms.

TABLE OF CONTENTS

List of Figures	iv
List of Tables	vi
List of Algorithms	vii
Abstract	viii
Εκτεταμένη Περίληψη	ix
1 Introduction	1
1.1 Scope	1
1.2 Structure	4
2 Preliminaries	5
2.1 Theory	5
2.1.1 Markov Chains	5
2.1.2 Absorbing Markov Chains	5
2.1.3 PageRank	7
2.1.4 Personalized PageRank	7
2.1.5 Relation Between PageRank and Absorbing Markov Chains. . .	8
2.2 Useful Notations	8
3 Fairness Aware PageRank Ranking	9
3.1 Definitions	9
3.2 Fairness Sensitive PageRank	11
3.2.1 The Algorithm	11
3.2.2 Optimizing Utility	11
3.2.3 Targeted Fairness Algorithm	12

3.3	Locally Fair PageRank	12
3.3.1	The Algorithms	13
3.3.2	Optimizing Utility	17
3.3.3	Targeted Fair Local Algorithms	17
3.4	A post processing approach	18
3.4.1	The Post Processing Algorithm	18
3.4.2	Targeted Fairness	20
4	PageRank Fair Recommendations	21
4.1	Theory	21
4.1.1	Impact of Adding an Edge	21
4.1.2	Impact of Adding a Set of Edges to a Single Source Node	25
4.1.3	Efficient Computation of Red Personalized PageRank and Per- sonalized PageRank	25
4.1.4	Fair Important Edges in a Network	26
5	Experimental Evaluation	27
5.1	Dataset Description	27
5.2	Fairness Aware PageRank Ranking	27
5.3	PageRank Fairness Aware Recommendations	34
5.3.1	Classic Recommendation Policies	36
5.3.2	Fair Recommendation Policies	38
5.3.3	Target Nodes Analysis	40
5.3.4	Homophily and Minority Size	43
5.3.5	Batch vs Online Gain	44
5.3.6	Fairness, Accepted Probability Correlation	45
6	Related Work	49
	Bibliography	51
A	Proofs	55
A.1	Fairness Aware PageRank Ranking	55
A.2	PageRankFair Recommendations	56

B Experiment Evaluation	57
B.1 Fairness Aware PageRank Ranking	57
B.1.1 Reproducibility	57
B.2 PageRankFair Recommendations	59

LIST OF FIGURES

5.1	Fairness of the PageRank algorithm with the size of the protected group for varying homophily.	28
5.2	Personalized pageranks starting from each of the blue nodes: histogram of the fraction of the personalized weight. The x-axis corresponds to weights fractions (blue, red and all (black bar)) and the y-axis to the percentage of the blues nodes with the corresponding fraction. The majority of nodes allocate the larger fraction of the personalized weight to blue nodes, thus being highly unfair to the opposite group.	29
5.3	Fairness sensitive Pagerank for $\phi = r$ and $\phi = 0.5$	30
5.4	Locally fair Pagerank algorithms for $\phi = 0.5$	31
5.5	Targeted fair PageRank algorithms and the optimal post-processing redistribution for $\phi = 0.5$. In each dataset, the target set S includes the 10% of nodes with the highest PageRank weights.	32
5.6	In neighborhood fairness for the nodes with the maximum loss and gain for each algorithm.	33
5.7	Visualization for $\phi = 0.5$	34
5.8	Visualization for $\phi = 0.5$	35
5.9	Impact on Fairness - Known Recommenders - Random Source Nodes.	37
5.10	Impact on Fairness - Known Recommenders - Red Source Nodes.	38
5.11	Impact on Fairness - Known Recommenders - Blue Source Nodes.	38
5.12	Impact on Fairness - Fair Recommenders - Random Source Nodes.	39
5.13	Impact on Fairness - Fair Recommenders - Red Source Nodes.	39
5.14	Impact on Fairness - Fair Recommenders - Blue Source Nodes.	40
5.15	Average Acceptance Probability - Random Source Nodes.	41
5.16	Average Acceptance Probability - Red Source Nodes.	41
5.17	Average Acceptance Probability - Blue Source Nodes.	42

5.18 Red PageRank ratio at top k nodes by PageRank - Blogs network. . . .	42
5.19 Red PageRank ratio at top k by nodes PageRank - Twitter network. . .	43
5.20 Cutting Point for Selecting Nodes.	44
5.21 Cutting Point for Selecting Nodes	45
5.22 Red PageRank ratio to different same group preference probability for sizes 0.1, 0.3, 0.5 after 5 and 10 link additions for random sources. . .	47
5.23 Fairness Impact for Batch and Online Fair Policy - Blogs.	48
5.24 Fairness Impact for Batch and Online Fair Policy - Twitter.	48
5.25 Recommendation Score - Fair Score Correlation.	48
 B.1 Cutting Point for Selecting Nodes.	 59
B.2 Cutting Point for Selecting Nodes.	61
B.3 Cutting Point for Selecting Nodes.	61
B.4 Cutting Point for Selecting Nodes.	61

LIST OF TABLES

3.1	Real dataset characteristics. r, b relative size of protected and unprotected group, respectively; p_R, p_B pagerank assigned to the red and blue group respectively	18
5.1	Utility loss with respect to optimal utility ($\frac{LFPR_X}{OPTIMAL}$, for $\phi = 0.5$) . . .	30
5.2	Number of Total Unique Target Nodes by Policy	41
5.3	Target Quality Features in Blogs Network.	43
5.4	Target Quality Features in Twitter Network.	43
B.1	Real dataset characteristics. r, b relative size of protected and unprotected group, respectively; p_R, p_B pagerank assigned to the red and blue group respectively	57
B.2	Utility loss with respect to optimal utility ($\frac{LFPR_X}{OPTIMAL}$)	58
B.3	Target Quality Features in Blogs - Red Source Nodes.	60
B.4	Target Quality Features in Blogs - Blue Source Nodes.	60
B.5	Target Quality Features in Twitter - Red Source Nodes.	60
B.6	Target Quality Features in Twitter - Blue Source Nodes.	60

LIST OF ALGORITHMS

3.1	Optimal Redistribution Algorithm	19
-----	--	----

ABSTRACT

Sotirios Tsioutsoulouklis, M.Sc. in Data and Computer Systems Engineering, Department of Computer Science and Engineering, School of Engineering, University of Ioannina, Greece, September 2020.

Fairness Aware Ranking & Recommendations in Networks.

Advisor: Evaggelia Pitoura, Professor.

Algorithmic fairness has attracted significant attention in the past years. Surprisingly, there is little work on fairness in networks. In this work, we consider fairness for link analysis algorithms and in particular for the celebrated PageRank algorithm. We provide definitions for fairness, and propose two approaches for achieving fairness. Furthermore, we explore how a recommendation system can affect the fairness of a network. We define objective for a fair recommender and we propose two recommendation policies in this direction. We present experiments with real and synthetic graphs that examine the fairness of PageRank, demonstrate qualitatively and quantitatively the properties of our fair algorithms and evaluate the impact of the different recommendation systems.

ΕΚΤΕΤΑΜΕΝΗ ΠΕΡΙΛΗΨΗ

Σωτήριος Τσιουτσιουλικλής, Δ.Μ.Σ. στη Μηχανική Δεδομένων και Υπολογιστικών Συστημάτων, Τμήμα Μηχανικών Η/Υ και Πληροφορικής, Πολυτεχνική Σχολή, Πανεπιστήμιο Ιωαννίνων, Σεπτέμβριος 2020.

Τεχνικές Κατάταξης και Πρότασης Νέων Σχέσεων σε Δίκτυα, Ενάντια στις Διακρίσεις.

Επιβλέπων: Ευαγγελία Πιτουρά, Καθηγήτρια.

Στην εποχή μας, λόγω του συνεχούς αυξανόμενου όγκου των δεδομένων προς επεξεργασία, χρησιμοποιούνται καθημερινά συστήματα και αλγόριθμοι για την ολοκλήρωση διάφορων διαδικασιών που μέχρι πρόσφατα διεξάγονταν από ανθρώπους. Συνήθεις διαδικασίες τέτοιων αλγορίθμων είναι η κατάταξη και η κατηγοριοποίηση των δεδομένων. Η εφαρμογή τέτοιων αλγορίθμων σε διαδικασίες που σχετίζονται με ανθρώπους (π.χ. 10 καλύτεροι ερευνητές για το 2020) είχαν ως αποτέλεσμα την εμφάνιση του ζητήματος των άκριτων διακρίσεων διαφόρων μορφών (π.χ. φυλετικές διακρίσεις) και της άνισης μεταχείρισης ανθρώπων από αλγορίθμους. Παρ' ότι το φαινόμενο έχει απασχολήσει την ερευνητική κοινότητα σε διάφορες κατηγορίες αλγορίθμων, όπως αυτών της μηχανικής μάθησης, και τα δίκτυα χρησιμοποιούντε στη μοντελοποίηση πληθώρας καθημερινών καταστάσεων και προβλημάτων, υπάρχει ελάχιστη δραστηριότητα προς αυτή τη κατεύθυνση στον τομέα των αλγορίθμων δικτύων.

Σε αυτή την εργασία επιχειρούμε μια πρώτη προσέγγιση στη καταπολέμιση των διακρίσεων σε αλγορίθμους που δρουν σε δίκτυα. Αρχικά, ορίζουμε τις έννοιες της δικαιοσύνης και του δίκαιου αλγορίθμου για δίκτυα. Επικεντρωνόμαστε στον δημοφιλή αλγόριθμο PageRank (Αν και η ανάλυση και οι αλγόριθμοι μπορούν να επεκταθούν κατά φυσικό τρόπο σε διάφορους άλλους αλγορίθμους για δίκτυα) και σε δυαδικά προστατευόμενα χαρακτηριστικά (π.χ. άντρας - γυναίκα), μελετάμε τις

ιδιότητες του δικτύου που το κάνουν άδικο και προτείνουμε διαφορετικές προσεγγίσεις προς τη παραγωγή ενός δίκαιου αποτελέσματος διατηρώντας παράλληλα εκείνα τα χαρακτηριστικά του αρχικού αλγορίθμου που τον ξεχωρίζουν και του προσδίδουν ιδιαίτερη αξία. Η πρώτη προσέγγιση χρησιμοποιεί τον διάνυσμα "άλματος" του PageRank για την επίτευξη ενός δίκαιου αποτελέσματος, ενώ η δεύτερη επιχειρεί μέσω της ατομικής συμπεριφοράς κάθε κόμβου αναγκάζοντας τον, κατά κάποιον τρόπο, να λειτουργήσει δίκαια. Επίσης, αξιολογούμε τους διαφορετικούς αλγορίθμους βάση της αλλαγής που φέρνουν σε σύγκριση με τον PageRank και τη χρησιμότητα τους. Οι αλγόριθμοι που προτείνουμε κλιμακώνουν αποδοτικά σε δεδομένα ευρείας κλίμακας.

Στη συνέχεια εξετάζουμε την επιρροή των συστημάτων συστάσεων συνδέσμων στη δικαιοσύνη ενός δικτύου. Παρατηρούμε ότι τα έως τώρα συστήματα συστάσεων δεν επιρεάζουν το δίκτυο σε αυτή τη παράμετρο, παρά διατηρούν την αρχική κατάσταση. Προτείνουμε ένα σύστημα συστάσεων που επιτυγχάνει την ανάδειξη και προβολή της αδικημένης/προστατευόμενης κατηγορίας στο δίκτυο με εξαιρετικά αποτελέσματα, θυσιάζοντας όμως τη ποιότητα των συστάσεων. Διατηρούμε το σκορ που παράγεται από το σύστημα αυτό και το εφραμόζουμε σε μια υβριδική μορφή σε συνδιασμό με ένα υπάρχον σύστημα συστάσεων. Για την πειραματική αξιολόγηση του συστήματος χρησιμοποιούμε ένα σύστημα συστάσεων βασισμένο σε embeddings προερχόμενα από τον node2vec αλγόριθμο και παρατηρούμε ότι το υβριδικό σύστημα ισορροπεί με ικανοποιητικό τρόπο τους δύο αντικειμενικούς στόχους μας (ανάδειξη της αδικημένης κατηγορίας και διατήρηση ποιοτικών συστάσεων). Επιπλέον, εξετάζουμε σε συνθετικά δίκτυα την συμπεριφορά των διαφόρων συστημάτων για διαφορετικές παραμέτρους και βλέπουμε ότι το προτεινόμενο σύστημα δεν επιρεάζεται από τα χαρακτηριστικά του δικτύου και συνεχίζει να έχει όμοια αποτελέσματα. Τέλος, μελετάμε τα ποιοτικά χαρακτηριστικά των συστάσεων όλων των αλγορίθμων και προσπαθούμε να εξηγήσουμε το σύστημα συστάσεων μέσα από απλά χαρακτηριστικά των προτεινόμενων συστάσεων.

CHAPTER 1

INTRODUCTION

1.1 Scope

1.2 Structure

1.1 Scope

Today, algorithmic systems driven by large amounts of data are increasingly being used in all aspects of life. Often, such systems are being used to assist, or, even replace human decision making. This increased dependence on algorithms has given rise to the field of algorithmic fairness, where the goal is to ensure that algorithms do not exhibit biases towards specific individuals, or groups of users (see e.g., [1] for a survey). We also live in a connected world where networks, be it, social, communication, interaction, or cooperation networks, play a central role. However, surprisingly, fairness in networks has received less attention.

Link analysis algorithms, such as Pagerank [2], HITS [3], or SALSA [4], take a graph as input and use the structure of the graph to determine the relative importance of its nodes. The output of the algorithms is a numerical weight for each node that reflects its importance. The weights are used to produce an ordering of the nodes and as input features in a variety of machine learning algorithms including classification [5], and search result ranking [2]. In this work, we focus on the Pagerank algorithm [2]. Pagerank performs a random walk on the input graph, and ranks the nodes according to the stationary probability of this walk. At every step, the random

walk restarts with probability c . The restart node is selected according to a “jump” distribution vector \mathbf{v} . Since its introduction in the Google search engine, Pagerank has been the cornerstone algorithm in several applications (see, e.g., [6]).

As in previous research, we view fairness as lack of discrimination against a protected group defined by the value of a sensitive attribute, such as, gender, or race [4]. We operationalize this view by saying that a link analysis algorithm is ϕ -fair, if the fraction of the total weight allocated to the members of the protected group is ϕ . The value of ϕ is a parameter that can be used to implement different fairness policies. In the simplest case, ϕ is set equal to the fraction of the protected nodes in the graph, asking that these nodes have a share in the weights proportional to their share in the population. We also consider *targeted* fairness, where we focus on a specific subset of nodes to which we want to allocate weight in a fair manner.

We revisit Pagerank through the lens of our fairness definition, and we consider the problem of defining Pagerank variants that are fair. We also define the *utility loss* of a fair algorithm as the difference between its output and the output of the Pagerank algorithm, and we pose the problem of achieving fairness while minimizing utility. We consider two approaches for achieving fairness. Our first approach, the *fairness-sensitive* Pagerank algorithm, exploits the jump vector \mathbf{v} . There has been a lot of work on modifying the jump vector to obtain variants of pagerank biased towards a specific set of nodes, for example, in personalized pagerank, all jump probability is assigned to a single node, while in topic-sensitive pagerank, the probability is assigned to nodes of a specific topic [7]. In this thesis, we take the novel approach of using the jump vector to achieve ϕ -fairness. We determine the conditions under which this is feasible and formulate the problem of finding the jump vector that achieves ϕ -fairness while minimizing utility loss from the original PageRank as a convex optimization problem.

Our second approach takes a microscopic view by looking at the behavior of each individual node in the graph. Implicitly, a link analysis algorithm assumes that links in the graph correspond to endorsements between the nodes. Therefore, we can view each node, as an agent that *endorses* (or *votes for*) the nodes that it links to. The link analysis algorithm defines a process that takes these individual actions of the nodes and transforms them into a global weighting of the nodes. To this end, we introduce, the *locally fair PageRank algorithms*, where each individual node acts fairly by distributing its own PageRank to the protected and non-protected groups according to the fairness ratio ϕ . Local fairness defines a dynamic process that can be viewed as

a *fair random walk*, where at each step of the random walk (not only at convergence), the probability of being at a node of the protected group is ϕ .

In our first locally fair PageRank algorithm, termed the *neighborhood locally fair* PageRank algorithm, each node distributes its PageRank fairly among its immediate neighbors, allocating a fraction ϕ to the neighbors in the protected group, and $1 - \phi$ to the neighbors in the non-protected group. Or, in random walk terms, at each node the probability of transitioning to a neighbor in the protected group is ϕ and the probability of transitioning to a non-protected neighbor is $1 - \phi$. The *residual-based locally fair* pagerank algorithms generalizes this idea. Consider a node i that has less neighbors in the protected group than ϕ . The node distributes an equal portion of its pagerank to each of its neighbors and a residual portion $\delta(i)$ to members in the protected group but not necessarily in its neighborhood. Or, in random walk terms, at each node i , the probability of transitioning to a neighbor is $1 - \delta(i)$ and the probability of transitioning to a node in the protected group is $\delta(i)$. The residual is allocated based on a *residual redistribution policy*, which allows us to control the fairness policy. In this thesis, we use the residual redistribution policy to minimize the utility loss.

Finally, we present a post-processing approach that given the output of a link analysis algorithm, it redistributes the weights so as to attain fairness. This gives us a lower bound on the utility loss.

We study the fairness of the original pagerank in both real and synthetic networks. We also evaluate quantitatively and qualitatively the output of our fairness-sensitive algorithms. The weights produced by the neighborhood locally fair Pagerank tend to promote protected nodes lying on the boundaries of the two groups especially in homophilic networks, while the fairness-sensitive Pagerank tends to jump to protected nodes especially when the requested ϕ is large.

Besides that, previous research has considered algorithms that weight nodes according to their degree, and found biases that arise as a network evolves [8, 9]. Link recommendation systems is known that can affect the evolution of a network. The first objective of a link recommendation system is to speed up the physical evolution of the network as this would have been done without any external intervention as this assumed to be more pleasant for the users. However, there are cases where changing a feature is important and recommendation systems are used for this purpose [10, 11]. In the context of fairness, as it is defined in the first part, we see that existing

recommenders preserve the initial bias of the network. We explore the possibilities that recommendations systems have and we show that they can be used to improve the fairness of a network.

At first, we call a recommender fair if the recommendations it makes improve the cumulative PageRank of the protected group. To do that we extend previous analysis that have been done in the effect of a new edge in PageRank[?] detect the edges that

In summary, in this thesis we make the following contributions:

- We initiate a study of fairness in link analysis. To the best of our knowledge, we are the first to consider fairness in this problem.
- We propose the fairness-sensitive Pagerank algorithm that modifies the jump vector so as to attain fairness and the locally fair Pagerank algorithms that guarantee that individually each node behaves in a fair manner
- We formulate optimization problems for finding the algorithms that minimizes the utility loss and estimate a lower bound for the optimal utility loss by post-processing the output of Pagerank
- We perform experiments on several datasets. Our experiments demonstrate qualitatively and quantitatively the properties of the fair Pagerank algorithms.

The remainder of this thesis is structured as follows. In Section 3.1, we define fair link analysis, while in Sections 3.2 and 3.3, we introduce the fairness sensitive and the locally fair PageRank algorithms. In Section 3.4, we present a post-processing approach for achieving fairness, while the results of our experimental evaluation are presented in Section ???. Finally, we compare our work with related research in Section ??? and offer our conclusions in Section ???.

1.2 Structure

CHAPTER 2

PRELIMINARIES

2.1 Theory

2.2 Useful Notations

2.1 Theory

2.1.1 Markov Chains

Markov chains are discrete stochastic processes on countable state spaces where the probability to move from one state to another depends only upon the present state. We call a state transient if the probability to leave from it is greater than 0 and absorbing otherwise. When the probability to move from every state to any other state is greater than 0 the Markov chain is called irreducible. Moreover, when for all states the gcd of the number of steps that is possible to go back to the current state is equal to 1 the Markov chain is called aperiodic. We call ergodic a Markov chain that is both irreducible and aperiodic. These stochastic processes have what we call as "stationary distribution". This means the existence of a distribution over the states, which describes the probability to visit each state, with the property that stays stable through time. We can express this as $\exists \mathbf{p} : \mathbf{p}^t = \mathbf{p}^t \cdot \mathbf{P} \wedge \mathbf{p}^T \cdot \mathbf{e} = 1, \mathbf{p} \in \mathbb{R}^n$.

2.1.2 Absorbing Markov Chains

An interesting category of stochastic processes are the absorbing Markov chains. An absorbing Markov chain is a Markov chain that has one or more absorbing states and

it is possible from every transient state to reach at least an absorbing one. We will now present the basic theory of absorbing Markov chains and a common modification of a Markov chain to an absorbing one, which provides us with the tools to prove interesting properties.

Use different \mathbf{P} notation for the pagerank and the absorbing Markov chains

As it is clear by now every state in an absorbing Markov chain is either transient or absorbing. If we consider that the transient states are the first k states and the absorbing ones the last $n - k$, $n = |V|$, then we can write the transition matrix in its "canonical form":

$$\mathbf{P} = \begin{bmatrix} \mathbf{P}_t & \mathbf{R} \\ \mathbf{0} & \mathbf{I} \end{bmatrix} \quad (2.1)$$

where \mathbf{P}_t is the transition matrix from transient to transient states, \mathbf{R} is the transition matrix from transient to absorbing states, $\mathbf{0}$ is the zero matrix denoting the probabilities from absorbing to transient states and \mathbf{I} is the identity matrix denoting the probabilities from absorbing to absorbing states.

Transition probabilities in m steps are given by \mathbf{P}^m , from 2.1 we get:

$$\mathbf{P}^m = \begin{bmatrix} \mathbf{P}_t^m & \mathbf{R}^* \\ \mathbf{0} & \mathbf{I} \end{bmatrix} \quad (2.2)$$

Since there is a probability to reach an absorbing state from any transient state (not necessarily in 1 step) is greater than 0 we have that $\lim_{m \rightarrow \infty} \mathbf{P}_{ij}^m = 0$, $\forall i, j \in \{i | i \in V \wedge i \text{ is a transient state}\}$. Also we have that $\mathbf{P}_{ii}^m = 1$, $\forall i \in \{i | i \in V \wedge i \text{ is an absorbing state}\}$. These implies that the $\lim_{m \rightarrow \infty} \mathbf{P}^m$ exists and it holds:

$$\lim_{m \rightarrow \infty} \mathbf{P}^m = \begin{bmatrix} \mathbf{0} & \mathbf{B} \\ \mathbf{0} & \mathbf{I} \end{bmatrix} \quad (2.3)$$

where \mathbf{B}_{ij} denotes the probability that has a random walk to gets absorbed from the transient state $k + j$ starting from state i , $1 \leq i \leq k, 1 \leq j \leq (n - k)$

We can see that the i, j element of matrix $\mathbf{N} = \mathbf{I} + \mathbf{P}_t^1 + \mathbf{P}_t^2 + \dots$ (the series converges and the matrix \mathbf{N} is well defined) gives us the expected number of visits to state j if we start the random walk from state i until we get absorbed by an absorbing state. Matrix \mathbf{N} is called the "fundamental matrix" of an absorbing Markov chain. Now, if we multiply both sides of equation of matrix \mathbf{N} by $(\mathbf{I} - \mathbf{P}_t)$ we get:

$$(\mathbf{I} - \mathbf{P}_t) \cdot \mathbf{N} = \mathbf{I} \quad (2.4)$$

which implies that $(\mathbf{I} - \mathbf{P}_t)$ is the invers of the matrix \mathbf{N} .

Knowing the fundamental matrix we are able now to calculate the absorbing probabilities for the transient states denoted as \mathbf{B} in the canonical form. This because the probability to be absorbed from state j starting from the transient state i is equal with the expected times I visit each state multiplied by the probability to be absorbed in the next step from that state, or more formally:

$$\mathbf{B} = \mathbf{NR} \quad (2.5)$$

\mathbf{R} as defined in the canonical form.

2.1.3 PageRank

The PageRank algorithm is the best-known link analysis algorithm, popularized by its application in the Google search engine. The scoring vector of the algorithm is the stationary distribution of a random walk on the graph G . We will use \mathbf{p} to denote this probability vector (which is the same as the scoring vector \mathbf{w}).

More formal, let $G = (V, E)$ be a graph where $V = \{0, 1, \dots, n\}$ is the set of nodes and $E \subseteq V \times V$ is the set of edges exist in the graph. PageRank is a link analysis algorithm $\mathcal{G}^n \rightarrow \mathbb{R}^n$, $\mathcal{G}^n :=$ the set of all graphs of size n , that defines a stochastic process based on the graph structure - and particular an ergodic Markov chain.

In PageRank the nodes are the states of the stochastic process and the transition probabilities are defined by the formula $\mathbf{P}_G = (1-a) \cdot \mathbf{P} + a \cdot \mathbf{e}\mathbf{v}^T$, $0 < a < 1$, where \mathbf{P}_G is the transition matrix defined by PageRank and \mathbf{P} is the transition matrix derived from the adjacency matrix of the graph with the difference that if a node has 0 out neighbors (it is an absorbing node) then we consider all nodes to be its neighbors, a is a factor known as "jump coefficient", \mathbf{v} is a probability vector (commonly the uniform one) known as "jump vector" and \mathbf{e} is the vector with 1 in all of its coordinates. PageRank vector can also be expressed as

$$\mathbf{p}^t = (1 - a)\mathbf{p}^t \cdot \mathbf{P} + a\mathbf{v}^t \quad (2.6)$$

2.1.4 Personalized PageRank

PageRank is a way of evaluating the nodes and assign to them a score that represents their importance in the network. A common problem is when we want to evaluate

the nodes not by their importance in the network but by their importance in the network for a specific node $u \in V$ (or for a specific set of nodes $U \subset V$)[Sotiris: Make text and formula generilized to sets]. In this case the PageRank is called personalized and this can be achieved by setting the jump vector

$$\mathbf{v}(i) = \begin{cases} 1, i = u \\ 0, otherwise \end{cases}$$

From 2.6 we can take that

$$\mathbf{p}^t = a\mathbf{v}^t \cdot \mathbf{Q}, \quad \mathbf{Q} = [\mathbf{I} - (1 - a) \cdot \mathbf{P}]^{-1} \quad (2.7)$$

The last relation implies that the PageRank of each node is the average personalized PageRank of that node, for all the nodes in the network.

2.1.5 Relation Between PageRank and Absorbing Markov Chains.

Last we must point out that for an ergodic Markov chained defined from the PageRank algorithm we can define a new absorbing Markov chain by adding some absorbing random nodes and by connecting each of the pre existing nodes in one of the new absorbing nodes with probability equal to the jump probability c . Then, the transition probabilities between transient states are defined as $\mathbf{P}_t = (1 - c)\mathbf{P}$ and $\mathbf{R}_{ij} = c$, if the transient state i has been connected to the new absorbing state $k + j$ and 0 otherwise. For this pair of Markov Chains, from equations 2.7, 2.4 it holds that the fundamental matrix of the new absorbing Markov chain is equal with the matrix of personalized PageRanks \mathbf{Q} .

$$\mathbf{N} = \mathbf{Q} \quad (2.8)$$

2.2 Useful Notations

Add notations

CHAPTER 3

FAIRNESS AWARE PAGERANK RANKING

3.1 Definitions

3.2 Fairness Sensitive PageRank

3.3 Locally Fair PageRank

3.4 A post processing approach

3.1 Definitions

A link analysis algorithm can be seen as a function $A : \mathcal{G}^n \rightarrow \mathbb{R}^n$ from the set \mathcal{G}^n of all graphs of size n to the real vectors of size n . The function takes as input a graph $G = (V, E)$ (directed, or undirected) of size n , and produces a vector \mathbf{w} of size n , which assigns a weight w_v to each node v in the graph. This weight defines the importance of the node in the graph G , and it depends on the graph structure. The best known link analysis algorithm is PageRank, which we consider in this paper.

Given a graph $G = (V, E)$, we assume that there exists a subset of nodes that define a *protected group*. This group may be defined based on a protected attribute of the nodes in the graph, such as race or gender. In this paper, we consider two types of nodes, the groups R and B of red and blue nodes, and we assume that R is the protected group. We denote with $r = \frac{|R|}{n}$, and $b = \frac{|B|}{n}$, the fraction of nodes that belong to the red and blue group respectively.

We will say that a link analysis algorithm is *fair*, if it assigns weights to each group according to a specified ratio ϕ . Ratio ϕ may be specified so as to implement specific

affirmative action policies, or other fairness enhancing interventions. For example, ϕ may be set in accordance to the 80 percent rule advocated by the US Equal Employment Opportunity Commission (EEOC), or some other formulation of disparate impact [12].

Definition 3.1 (Fair link analysis). A link analysis algorithm $A : \mathcal{G}^n \rightarrow \mathbb{R}^n$ is ϕ -fair on graph G , if for the output $\mathbf{w} = A(G)$, it holds that: $\frac{\sum_{v \in R} w_v}{\sum_{v \in V} w_v} = \phi$, where $R \subset V$ is the protected set of nodes.

For instance by setting $\phi = r$, we ask for a fair link analysis algorithm that assigns weights proportionally to the sizes of the two groups. In this case, fairness is analogous to demographic parity, i.e., the requirement that the demographics of those receiving a positive outcome are identical to the demographics of the population as a whole [13]. It is easy to show (see Appendix) that in this case the weights produced by the fair link analysis are such that the average weight of the red nodes is the same with the average weight of the blue nodes.

We define the following problem:

Problem 1. *Given a value ϕ , a graph G , and a link analysis algorithm A , design a link analysis algorithm A_F that is ϕ -fair on graph G .*

Note that the fair variant A_F will necessarily change the original weights of algorithm A , incurring some loss in *utility*. We quantify the *utility loss* using the sum of squares loss function $L(A, A_F) = \|A(G) - A_F(G)\|^2$. We then consider the problem of designing a fair algorithm that minimizes utility loss.

Problem 2. *Given a value ϕ , a graph G , and a link analysis algorithm A , design a link analysis algorithm A_F that is ϕ -fair on graph G , such that the utility loss $L(A(G), A_F(G))$ is minimized.*

Finally, we consider an extension of the fairness definition that asks for a fair distribution of weights among a specific set of nodes S that is given as input. We assume that the set S is selected such that it contains nodes from both groups R and B .

Definition 3.2 (Targeted Fair link analysis). A link analysis algorithm $A : \mathcal{G}^n \rightarrow \mathbb{R}^n$ is targeted ϕ -fair on graph $G = (V, E)$ for a set of nodes $S \subset V$, if for the output $\mathbf{w} = A(G)$, it holds that $\frac{\sum_{v \in S \cap R} w_v}{\sum_{v \in S} w_v} = \phi$, where $R \subset V$ is the protected set of nodes.

In this paper, we consider the PageRank link analysis algorithm.

3.2 Fairness Sensitive PageRank

Our first algorithm achieves fairness by keeping the transition matrix fixed and changing the jump vector \mathbf{v} so as to meet the fairness criterion.

3.2.1 The Algorithm

First, we note that that pagerank vector \mathbf{p} can be written as linear function of the jump vector \mathbf{v} . Solving Equation (??) for \mathbf{p} , using column vector notation, we have that $\mathbf{p} = \mathbf{Q}\mathbf{v}$, where

$$\mathbf{Q} = \gamma ([\mathbf{I} - (1 - \gamma)\mathbf{P}]^{-1})^T$$

Let \mathbf{p}_R denote the pagerank mass that is allocated to the nodes of the protected category. We have that

$$\mathbf{p}_R = \left(\sum_{i \in R} \mathbf{Q}\mathbf{v} \right) [i] = \left(\sum_{i \in R} \mathbf{Q}_i^T \right) \mathbf{v} = \mathbf{Q}_R^T \mathbf{v}$$

where \mathbf{Q}_i^T is the i -th row of matrix \mathbf{Q} , and \mathbf{Q}_R^T is the vector that is the sum of the rows in the set R . In order for the algorithm to be fair, we need $\mathbf{p}_R = \phi$. Our goal is to find a vector \mathbf{v} such that $\mathbf{Q}_R^T \mathbf{v} = \phi$.

Does such a vector always exist? We can prove the following:

Lemma 3.1. *Given the vector \mathbf{Q}_R^T , there exists a vector \mathbf{v} such that $\mathbf{Q}_R^T \mathbf{v} = \phi$, if and only if, there exist entries i, j in \mathbf{Q}_R^T , where $\mathbf{Q}_R^T(i) \leq \phi$ and $\mathbf{Q}_R^T(j) \geq \phi$*

Proof. We have that $\mathbf{p}_R = \sum_{j=1}^N \mathbf{Q}_R^T(j)v_j$, that is, \mathbf{p}_R is the weighted average of the values $\mathbf{Q}_R^T(j)$, with weights v_j , where $0 \leq v_j \leq 1$. Since $\mathbf{Q}_R^T \mathbf{v} = \phi$, there must exist at least one entry i with $\mathbf{Q}_R^T(i) \leq \phi$, and one entry j $\mathbf{Q}_R^T(j) \geq \phi$. Conversely, if there exists two such entries i, j , then we can find values v_i and v_j , such that $v_i \mathbf{Q}_R^T(i) + v_j \mathbf{Q}_R^T(j) = \phi$ and $v_i + v_j = 1$. \square

3.2.2 Optimizing Utility

An implication of Lemma 3.1 is that, in most cases, there are multiple jump vectors that give a fair pagerank vector. We are interested in the solution that minimizes the utility loss.

We first consider the case were we want fairness over all nodes. To solve this problem we exploit the fact that the utility loss function $L(\mathbf{p}_v, \mathbf{p}_u) = \|\mathbf{p}_v - \mathbf{p}_u\|^2$ is

convex, and that we can express the fairness requirement as a linear function. We can then define the following convex optimization problem.

$$\begin{aligned}
& \underset{\mathbf{x}}{\text{minimize}} && \|\mathbf{Q}\mathbf{x} - \mathbf{p}_u\|^2 \\
& \text{subject to} && \mathbf{Q}_R^T \mathbf{x} = \phi \\
& && \sum_{i=1}^n \mathbf{x}_i = 1 \\
& && 0 \leq \mathbf{x}_i \leq 1, \ i = 1, \dots, n
\end{aligned}$$

This problem can be solved using standard convex optimization solvers.

3.2.3 Targeted Fairness Algorithm

We will now formulate a similar convex optimization problem for the targeted fairness problem. Let $\mathbf{Q}_S^T = \sum_{i \in S} \mathbf{Q}_i^T$ be the sum of rows of \mathbf{Q} for the nodes in S , and $\mathbf{Q}_{R|S}^T = \sum_{i \in S \cap R} \mathbf{Q}_i^T$ be the sum of rows of \mathbf{Q} for the R nodes in S . We define a convex optimization problem that is exactly the same as in Section 3.2.2, except for the fact that we replace the constraint $\mathbf{Q}_R^T \mathbf{x} = \phi$ with the constraint $\mathbf{Q}_{R|S}^T \mathbf{x} = \phi \mathbf{Q}_S^T \mathbf{x}$.

We can model specific cases by adding additional constraints. For example, let $T_k(\mathbf{w})$ denote the k nodes with the largest weights in vector \mathbf{w} , and let $S = T_k(\mathbf{p}_u)$, that is, the top- k nodes of the original Pagerank algorithm. We want fair redistribution of Pagerank among the nodes in S , but we also want these nodes to remain in the top- k position in the fair pagerank, that is, $T_k(\mathbf{p}) = T_k(\mathbf{p}_u)$. This requirement can be achieved by adding the constraint:

$$\mathbf{Q}_i^T \mathbf{x} \geq \mathbf{Q}_j^T \mathbf{x}, \ i \in T_k(\mathbf{p}), j \notin T_k(\mathbf{p}).$$

3.3 Locally Fair PageRank

The locally fair PageRank algorithms take a microscopic view, by asking that *each individual node* acts fairly, i.e., each node distributes its own pagerank to red and blue nodes fairly. In random walk terms, local fairness defines a dynamic process that can be viewed as a random walk that is fair, i.e., at each step, and not just at convergence, the probability of being at a node of the protected group is ϕ .

3.3.1 The Algorithms

In our first algorithm, the *neighborhood locally fair PageRank* algorithm, each node distributes its own pagerank fairly among the red and blue nodes in its neighbors. The *residual-based locally fair PageRank algorithms* generalize this idea, by again asking that each node allocates its pagerank fairly among the red and blue nodes, but not necessarily among the red and blue nodes in its own neighborhood. Finally, we show that the local PageRank algorithms are fair.

The neighborhood locally fair PageRank algorithm

We first consider a node that treats its neighbors fairly, that is, by allocating a fraction ϕ of its pagerank to its red neighbors and the remaining $1 - \phi$ fraction to its blue neighbors. In random walk terms, at each node the probability of transitioning to a red neighborhood is ϕ and the probability of transitioning to a blue neighborhood $1 - \phi$.

Specifically, we define the *neighborhood locally fair pagerank* (LFPR_N) \mathbf{p}_N as follows. Each node i splits the $\phi \mathbf{p}_N(i)$ portion of its pagerank value evenly among its red out-neighbors and the remaining $(1 - \phi) \mathbf{p}_N(i)$ portion of its pagerank evenly among its blue out-neighbors. Similarly, we use a modified “fair” jump vector \mathbf{v}_N with $\mathbf{v}_N[i] = \frac{\phi}{|R|}$, if $i \in R$, and $\mathbf{v}_N[i] = \frac{1-\phi}{|B|}$, if $i \in B$.

Let $out_R(i)$ and $out_B(i)$ be the number of edges directed from node i to red nodes and blue nodes respectively. We define \mathbf{P}_R as the normalized adjacency matrix that includes links to red nodes, or random jumps to red nodes if such links do not exist:

$$\mathbf{P}_R(i, j) = \begin{cases} \frac{1}{out_R(i)}, & \text{if } j \in R, out_R(i) \neq 0, \text{ and } (i, j) \in E \\ \frac{1}{|R|}, & \text{if } j \in R, \text{ and } out_R(i) = 0 \\ 0, & \text{otherwise} \end{cases}$$

\mathbf{P}_B is defined similarly. The transition matrix \mathbf{P}_N of the LFPR_N algorithm is:

$$\mathbf{P}_N = \phi \mathbf{P}_R + (1 - \phi) \mathbf{P}_B$$

and, the neighborhood locally-fair pagerank vector \mathbf{p}_N is defined as:

$$\mathbf{p}_N^T = (1 - \gamma) \mathbf{p}_N^T \mathbf{P}_N + \gamma \mathbf{v}_N^T$$

The neighborhood locally-fair pagerank value \mathbf{p}_N of a node is the stationary probability that a neighborhood-fair walker ends up at this node.

The residual-based locally fair PageRank algorithms

We consider an alternative fair behavior for individual nodes. Similarly to the LFPR_N algorithm, each node i acts fairly by respecting the ϕ ratio when distributing its own pagerank to red and blue nodes. However, now node i treats its neighbors the same, independently of their color and assigns to each of them the same portion of its pagerank. When a node is in a “biased” neighborhood, i.e., the ratio of its red neighbors is different than ϕ , to be fair, node i distributes the remaining portion of its pagerank to nodes in the underrepresented group. We call the remaining portion *residual* and denote it by $\delta(i)$. How $\delta(i)$ is distributed to the underrepresented group is determined by a *residual policy*.

Intuitively, this corresponds to a fair random walker that upon arriving at a node i , with probability $1-\delta(i)$ follows one of i ’s outlinks and with probability $\delta(i)$ jumps to one or more nodes in the locally underrepresented group.

We now describe the algorithm formally. We divide the nodes in V into two sets, L_R and L_B , based on the fraction of their red and blue neighbors. Set L_R includes the “blue-biased” nodes, that is, all nodes i such that $(1 - \phi) \text{out}_R(i) < \phi \text{out}_B(i)$, that is, the nodes for which the ratio of red nodes in their neighborhood is smaller than the required ϕ ratio. These are the nodes having a residual that needs to be distributed to red nodes. Analogously, L_B includes all “red-biased” nodes, that is, all nodes i such that $(1 - \phi) \text{out}_R(i) \geq \phi \text{out}_B(i)$.

Let us first consider a node i in L_R . Each neighbor of i gets the same portion of i ’s pagerank, let $\rho_R(i)$ be this portion. To attain the ϕ ratio, the residual $\delta_R(i)$ of i ’s pagerank goes to the red nodes. Portions $\rho_R(i)$ and $\delta_R(i)$ must be such that:

$$(1 - \phi) (\text{out}_R(i) \rho_R(i) + \delta_R(i)) = \phi (\text{out}_B(i) \rho_R(i)) \quad (3.1)$$

$$\text{out}_R(i) \rho_R(i) + \text{out}_B(i) \rho_R(i) + \delta_R(i) = 1 \quad (3.2)$$

From Equations (3.1) and (3.2), we get $\rho_R(i) = \frac{1-\phi}{\text{out}_B(i)}$ and the residual is $\delta_R(i) = \phi - \frac{(1-\phi) \text{out}_R(i)}{\text{out}_B(i)}$.

Analogously, for a node i in L_B , we get $\rho_B(i) = \frac{\phi}{\text{out}_R(i)}$ and a residual $\delta_B(i) = (1 - \phi) - \frac{\phi \text{out}_B(i)}{\text{out}_R(i)}$ that goes to the blue nodes.

Example. Consider a node i with 5 out-neighbors, 1 red and 4 blue, and let ϕ be 0.5.

This is a “blue-biased” node, that is a node in L_R . In the original PageRank algorithm, each of the 5 neighbors gets $1/5$ of i ’s pagerank, resulting in red nodes getting $1/5$ and blue nodes $4/5$ of i ’s pagerank, which is an unfair behavior for node i . With the residual algorithm, each of i ’s neighbors gets $\rho_R(i) = 1/8$ portion of i ’s Pagerank, resulting in red neighbors getting $1/8$ and blue neighbors $4/8$ of i ’s pagerank. The residual $\delta_B(i) = 3/8$ goes to nodes in the red group so as to attain the ϕ ratio and make i fair. Which of the nodes in the red group will get the residual is determined by the residual policy. In terms of the random walker interpretation, a random walker that arrives at i , with probability $5/8$ chooses one (any) of i ’s outlinks and with probability $3/8$ jumps to nodes in the red group.

The transition matrix \mathbf{P}_L is defined as

$$\mathbf{P}_L(i, j) = \begin{cases} \frac{1-\phi}{out_B(i)}, & \text{if } (i, j) \in E \text{ and } i \in L_R \\ \frac{\phi}{out_R(i)}, & \text{if } (i, j) \in E \text{ and } i \in L_B \\ 0, & \text{otherwise} \end{cases}$$

Let δ_R be the vector carrying the red residual, that is, $\delta_R[i] = \phi - \frac{(1-\phi)out_R(i)}{out_B(i)}$, if $i \in L_R$ and 0 otherwise. Similarly, let δ_B be the vector carrying the blue residual, that is, $\delta_B(i) = (1 - \phi) - \frac{\phi out_B(i)}{out_R(i)}$, if $i \notin L_B$ and 0 otherwise. We have a total red residual $\Delta_R = \mathbf{p}_L^T \delta_R$ and a total blue residual $\Delta_B = \mathbf{p}_L^T \delta_B$, where \mathbf{p}_L is the locally fair pagerank vector.

To express the residual distribution policy, we introduce two matrices, matrices \mathbf{X} and \mathbf{Y} , that capture the policy for distributing the residual to red and blue nodes respectively. Specifically, $\mathbf{X}[i, j]$ denotes the portion of the $\delta_R(i)$ of node $i \in L_R$ that goes to node $j \in R$ and $\mathbf{Y}[i, j]$ the portion of the $\delta_B(i)$ of node $i \in L_B$ that goes to node $j \in B$.

The locally-fair pagerank vector \mathbf{p}_L is defined as:

$$\mathbf{p}_L^T = (1 - \gamma)\mathbf{p}_L^T (\mathbf{P}_L + \mathbf{X} + \mathbf{Y}) + \gamma \mathbf{v}_N^T$$

Residual Distribution Policies. The \mathbf{X} and \mathbf{Y} allocation matrices allow us the flexibility to specify appropriate policies for distributing the residual. In particular, the following holds (proof in the Appendix):

Lemma 3.2. *The LFPR_N algorithm is a special case of the residual-based algorithm, with*

$$\mathbf{X}_N[i, j] = \begin{cases} \frac{1}{\text{out}_R(j)}, & \text{if } i \in R, j \in L_R, \text{ and } (j, i) \in E \\ 0 & \text{otherwise} \end{cases}$$

$$\mathbf{Y}_N[i, j] = \begin{cases} \frac{1}{\text{out}_B(j)}, & \text{if } i \in B, j \in L_B, \text{ and } (j, i) \in E \\ 0 & \text{otherwise} \end{cases}$$

We also consider residual policies where all nodes follow the same policy in distributing their residual. In this case, the residual policy is expressed through two (column) vectors \mathbf{x} and \mathbf{y} , with $\mathbf{x}[i]$ being the portion of Δ_R going to red node i , and $\mathbf{y}[i]$ the portions of Δ_B going to blue node i . In this case, we have:

$$\mathbf{p}_L^T = (1 - \gamma)\mathbf{p}_L^T (\mathbf{P}_L + \delta_R \mathbf{x}^T + \delta_B \mathbf{y}^T) + \gamma \mathbf{v}_N^T.$$

We define two locally fair PageRank algorithms based on two intuitive policies of distributing the residual, namely:

- (1) the *Uniform Locally Fair PageRank* (LFPR_U) algorithm, which distributes the residual uniformly. Specifically, for the LFPR_U algorithm, we define the vector \mathbf{x} , as $\mathbf{x}[i] = \frac{1}{|R|}$ if $i \in R$ and 0 otherwise, and the vector \mathbf{y} , as $\mathbf{y}[i] = \frac{1}{|B|}$, if $i \in B$ and 0 otherwise.
- (2) the *Proportional Locally Fair PageRank* (LFPR_P) algorithm, which distributes the residual proportionally based on the original pagerank weight $\mathbf{p}_u(i)$ of node i . Specifically, for the LFPR_P algorithm, we define the vector \mathbf{x} , as $\mathbf{x}[i] = \frac{\mathbf{p}[i]}{\sum_{i \in R} \mathbf{p}[i]}$, if $i \in R$ and 0 otherwise, and the vector \mathbf{y} , as $\mathbf{y}[i] = \frac{\mathbf{p}^{(i)}}{\sum_{i \in B} \mathbf{p}[i]}$, if $i \in B$ and 0, otherwise.

Fairness of the locally fair PageRank algorithms

In the locally fair pagerank algorithms, each node in the graph treats the red and blue nodes fairly by respecting the ϕ ratio. However, each node acts independently of the other nodes in the network. It is interesting to see how this microscopic view of fairness relates to our macroscopic view of link fairness.

We prove the following theorem.

Theorem 3.1. *The locally fair pagerank algorithms are fair.*

Proof. We must show that $\frac{\sum_{v \in R} \mathbf{p}_N(u)}{\sum_{v \in V} \mathbf{p}_N(u)} = \phi$. Since each node in the graph gives a portion ϕ of its pagerank to red nodes, we have

$$\sum_{v \in R} \mathbf{p}_N(u) = \sum_{v \in V} \phi \mathbf{p}_N(u)$$

which proves the theorem. \square

3.3.2 Optimizing Utility

We consider how to optimally distribute the residual so as to minimize the utility loss of the fair Pagerank. We denote this algorithm as LFPR_O. To this end, we define appropriate \mathbf{x} and \mathbf{y} residual distribution vectors by formulating an optimization problem.

We can write the vector \mathbf{p}_L as a function of the vectors \mathbf{x} and \mathbf{y} as follows:

$$\mathbf{p}_L^T(\mathbf{x}, \mathbf{y}) = \gamma \mathbf{v}^T [\mathbf{I} - (1 - \gamma)(\mathbf{P}_L + \delta_R \mathbf{x}^T + \delta_B \mathbf{y}^T)]^{-1}$$

We can now define the optimization problem of finding the vectors \mathbf{x} and \mathbf{y} that minimize the loss function $L(\mathbf{p}_L, \mathbf{p}_u) = \|\mathbf{p}_L(\mathbf{x}, \mathbf{y}) - \mathbf{p}_u\|^2$ subject to the constraint that the vectors \mathbf{x} and \mathbf{y} define a distribution over the nodes in R and B respectively.

We solve this optimization problem using gradient descent. We enforce the distribution constraints by adding a penalty term $\lambda ((\sum_{i=1}^n \mathbf{x}_i - 1)^2 + (\sum_{i=1}^n \mathbf{y}_i - 1)^2)$. We enforce the positivity constraints through proper bracketing at the line-search step.

Note that we can also formulate a convex optimization problem asking for the jump vector that minimizes utility loss, as in Section 3.2.2. In this case, since the transition matrix is fair, we just need to constrain the jump vector to obey the ϕ ratio.

3.3.3 Targeted Fair Local Algorithms

We show how to apply the local algorithms to the targeted fairness problem. Let S_R and S_B be the red and blue nodes in the set S respectively, and let I_S be the set of in-neighbors of S . The idea is that the nodes in I_S should distribute their PageRank to S_R and S_B fairly, such that the ratio of the portion that goes to nodes in S_R and the portion that goes to nodes in S_B is equal to $\frac{\phi}{1-\phi}$. We can implement the same redistribution policies as in the case of the neighborhood local and the residual-based local fair algorithms.

We also need the (global) jump vector \mathbf{v} to obey the ϕ ratio for the nodes in S . We can achieve this by redistributing the probability $|S|/n$ of the jump vector according to the ϕ ratio. Note that there is a variety of policies one could implement, depending on a specific objective. For example if we want to increase the weight of the nodes in S , we can make the jump vector allocate all probability to the nodes in S .

3.4 A post processing approach

We now consider a post processing approach in which we assume that we are given a weight vector $\mathbf{w} = A(G)$ of a link analysis algorithm A on graph G . The goal is to produce a new weight vector \mathbf{f} such that: (1) \mathbf{f} is fair, and (2) the utility loss $L(\mathbf{w}, \mathbf{f}) = \|\mathbf{w} - \mathbf{f}\|^2$ is minimized. The post-processing algorithm is agnostic to the fact that the weight vector \mathbf{w} is the result of a link analysis algorithm, much less of the specific link analysis algorithm (e.g., Pagerank). Therefore, the vector \mathbf{f} that minimizes the loss $L(\mathbf{w}, \mathbf{f})$ may not be attainable by any Pagerank algorithm.

Table 3.1: Real dataset characteristics. r , b relative size of protected and unprotected group, respectively; p_R , p_B pagerank assigned to the red and blue group respectively

Dataset	#nodes	#edges	Protected attribute	r	b	homophily	p_R	p_B
BOOKS	92	748	political (left)	0.47	0.53	0.06	0.47	0.53
BLOGS	1,222	19,089	political (left)	0.52	0.48	1.18	0.37	0.63
TWITTER	18,470	61,157	political (left)	0.61	0.39	0.04	0.57	0.43
DBLP2	13,015	79,972	gender (women)	0.17	0.83	0.87	0.16	0.84

3.4.1 The Post Processing Algorithm

Given the weight vector \mathbf{w} , let \mathbf{w}_R denote the weight vector for the nodes in R , and \mathbf{w}_B the weight vector for the nodes in B . We also use W_R to denote the total weight allocated to R , and W_B to denote the total weight allocated to B . We assume that \mathbf{w} has non-negative entries, and it is normalized so that its entries sum to 1. Without loss of generality assume that $W_R < \phi$. Let $\Delta = \phi - W_R$. To make the vector fair we need to distribute weight Δ to the nodes in R , and remove weight Δ from the nodes in B . It is easy to show that in order to minimize the loss, the optimal redistribution

Algorithm 3.1 Optimal Redistribution Algorithm

Input: Excess weight Δ , nodes B , weights \mathbf{w}_B

Output: Optimal weight vector \mathbf{f}_B

```
1:  $B_{NZ} \leftarrow \{x \in B : w_x > 0\}$ 
2:  $\delta = \Delta / |B_{NZ}|$ 
3:  $\beta = \min_{x \in B_{NZ}} w_x$ 
4: if  $\beta \geq \delta$  then
5:    $w_x = w_x - \delta$  for all  $x \in B_{NZ}$ 
6:
7:   return  $\mathbf{w}_B$ 
8: else
9:    $w_x = w_x - \beta$  for all  $x \in B_{NZ}$ 
10:   $\Delta = \Delta - |B_{NZ}| \beta$ 
11:
12:  return REDISTRIBUTE( $\Delta, B, \mathbf{w}_B$ )
13: end if
```

will remove weight $\Delta/|B|$ from all nodes in B and add $\Delta/|R|$ from all nodes in R . This follows from the fact that among all distribution vectors the one with the smallest length is the uniform one. Therefore, we obtain the following lower-bound for the loss:

$$\text{Loss}_{LB} = \frac{\Delta^2}{|R|} + \frac{\Delta^2}{|B|}$$

Note that this lower bound does not guarantee that the new vector \mathbf{f} has non-negative entries, thus it is not a valid weight vector. We now describe an optimal redistribution algorithm that ensures that when removing weight no entry becomes negative, while using the principle that whenever removing weight, the optimal way is to remove uniformly from all nodes. The pseudocode for the algorithm is shown in Algorithm 3.1.

Solve matter with procedure

The algorithm takes as input the value of excess weight Δ that needs to be removed, the set of nodes B from which we want to remove the weight, and the current weights \mathbf{w}_B of these nodes. First, it finds the subset of nodes B_{NZ} in B that have non-zero weight. If the minimum weight β among these nodes is at least $\Delta/|B_{NZ}|$,

then we can remove the weight uniformly without making the weights negative. The algorithm updates the weights and returns. Otherwise, we can remove at most β . The algorithm removes β from all nodes in B_{NZ} and makes a recursive call with the remaining excess weight $\Delta - |B_{NZ}|\beta$. Note that anytime we want to remove weight from a set of nodes, we remove it uniformly from all nodes, which guarantees optimality. The algorithm returns the updated weight vector \mathbf{f}_B for the nodes in B . We can now compute the optimal loss as

$$\text{Loss}_O = \frac{\Delta^2}{|R|} + \|\mathbf{f}_B - \mathbf{w}_B\|^2$$

3.4.2 Targeted Fairness

Computing algorithmically the optimal redistribution is harder in the targeted fairness case, since there are many different options in how we can redistribute weight. We can move weight between the nodes in S , or bring in weight from outside of S , or move weight out of S , or a combination of those. In Appendix A.1 we compute analytically a lower bound for the loss, which provides some intuition on how the weight is moved in different cases .

Finding the optimal redistribution vector can be formulated as a convex optimization problem:

$$\begin{aligned} & \underset{\mathbf{f}}{\text{minimize}} && \|\mathbf{f} - \mathbf{p}\|^2 \\ & \text{subject to} && \sum_{i \in S \cap R} \mathbf{f}_i = \phi \\ & && \sum_{i=1}^n \mathbf{f}_i = 1 \\ & && 0 \leq \mathbf{f}_i \leq 1, \ i = 1, \dots, n \end{aligned}$$

We use the solution of the optimization problem to compare the optimal redistribution with that achieved by the modified Pagerank algorithms.

CHAPTER 4

PAGERANK FAIR RECOMMENDATIONS

4.1 Theory

4.1 Theory

4.1.1 Impact of Adding an Edge

Lets assume an unweighted, directed graph $G = (V, E)$ where V is the set of all nodes and E is the set of all edges. And let $G' = (V, E' = E \cup \{(u, v)\})$ where $(u, v) \notin E$. We denote with $\mathbf{p}(R)$, $\mathbf{p}'(R)$ the ratio of the pagerank that goes to Red nodes in the graph G and G' respectively. We denote by k_u the out degree of node u

Now let \mathbf{Q} be the matrix where \mathbf{Q}_{ij} is the personilized PageRank of node i to node u and so $\mathbf{Q}_i(R)$ is the personilized PageRank of node i to all Red nodes.

For \mathbf{p} and \mathbf{p}' we can prove the following.

Theorem 4.1. *If $G' = (V, E' = E \cup \{u, v\})$ then:*

$$\mathbf{p}'(R) = \mathbf{p}(R) + \mathbf{p}_u \cdot \frac{\frac{(1-c)}{c} [\mathbf{Q}_v(R) - \frac{1}{k_u} \sum_{w \in E_u} \mathbf{Q}_w(R)]}{(k_u + 1) - \frac{(1-c)}{c} [\mathbf{Q}_{vu} - \frac{1}{k_u} \sum_{w \in E_u} \mathbf{Q}_{wu}]}$$

Proof. To prove this we first write the transition matrix \mathbf{P} of G' as a sum of the tranistion matrix \mathbf{P} of graph G and a rank one matrix \mathbf{D} . This is:

$$\mathbf{P}' = \mathbf{P} + \mathbf{D}, \quad \mathbf{D}_i = \begin{cases} 0, & i \neq u \\ (-1 + \frac{k_u}{k_u+1})\mathbf{P}_u + \frac{1}{k_u+1}\mathbf{e}_v \end{cases}$$

And then we exploit a fundamental lemma [14] that states that If G is nonsingular, H is of rank 1 and $G + H$ is nonsingular as well, then:

$$(G + H)^{-1} = G^{-1} - \frac{1}{1 + g} G^{-1} H G^{-1}, \quad g := \text{tr}(H G^{-1})$$

We also know

Add reference here

that $\mathbf{p} = \mathbf{v}\mathbf{Q}$ and [15] that $\mathbf{Q} = c \cdot \mathbf{N}$ where \mathbf{N} is the fundamental matrix of an absorbing Markov chain with transition matrix for transient states equal to $(1 - c)\mathbf{P}$, c is the jump probability of PageRank.

For $G = [\mathbf{I} - (1 - c) \mathbf{P}]$ and $H = -(1 - c) \cdot \mathbf{D}$, we have:

$$\begin{aligned} \mathbf{N}' &= \mathbf{N} - \frac{1}{1 + q} \mathbf{N}(-(1 - c)\mathbf{D}\mathbf{N}), \quad q := \text{tr}(-(1 - c)\mathbf{D}\mathbf{N}) \Rightarrow \\ \frac{1}{c}\mathbf{Q}' &= \frac{1}{c}\mathbf{Q} - \frac{1}{c^2} \frac{1}{1 + q} \mathbf{Q}(-(1 - c) \cdot \mathbf{D})\mathbf{Q}, \quad q := \text{tr}(-(1 - c)\mathbf{D}\frac{1}{c}\mathbf{Q}) \Rightarrow \end{aligned}$$

$$\frac{1}{c}\mathbf{Q}' = \mathbf{Q} + \frac{1}{c} \frac{(1 - c)}{1 - \frac{(1 - c)}{c}q} \mathbf{Q}\mathbf{D}\mathbf{Q}, \quad q := \text{tr}(\mathbf{D}\mathbf{Q}) \quad (4.1)$$

If we compute $\mathbf{D}\mathbf{Q}$, $\mathbf{Q}\mathbf{D}\mathbf{Q}$ we get:

$$\begin{aligned} \mathbf{D}\mathbf{Q}_{ij} &= \begin{cases} 0, & i \neq u \\ \frac{1}{k_u + 1} [\cdot \mathbf{Q}_{vj} - \frac{1}{k_u} \sum_{w \in E_u} \mathbf{Q}_{wj}], & i = u \end{cases} \\ \mathbf{Q}\mathbf{D}\mathbf{Q}_{ij} &= \frac{1}{k_u + 1} (\mathbf{Q}_{vu} - \frac{1}{k_u} \sum_{w \in E_u} \mathbf{Q}_{wj}) \end{aligned}$$

And from (1):

$$\mathbf{Q}'_{ij} = \mathbf{Q}_{ij} + \mathbf{Q}_{iu} \frac{\frac{(1 - c)}{c} (\mathbf{Q}_{vj} - \frac{1}{k_u} \sum_{w \in E_u} (\mathbf{Q}_{wj}))}{k_u + 1 - \frac{(1 - c)}{c} (\mathbf{Q}_{vu} - \frac{1}{k_u} \sum_{w \in E_u} (\mathbf{Q}_{wu}))}$$

Since $\mathbf{p} = \mathbf{v}\mathbf{Q}$ we have to compute

$$\mathbf{p}'(R) = \frac{1}{n} \sum_{i=1}^n \sum_{j \in R} \mathbf{Q}'_{ij}$$

to obtain the formula of the theorem. □

We will show now that the denominator of the fraction is always positive. To do this we define an absorbing Markov chain \bar{X} . We add two absorbing states $n+1$, $n+2$, from now on a_u , a_o and we connect a state u to state a_u and all other states to state a_o , all with probability c . We denote with \mathbf{B}_{i1} the absorbing probability of node i to node a_u and \mathbf{Q}_{iu} the personalized PageRank of node i for the node u . The following holds.

Lemma 4.1. *The absorption probability of state i to state a_u of the absorbing Markov chain \bar{X} is equal with the personalized PageRank of node i to node u .*

$$\mathbf{Q}_{iu} = \mathbf{B}_{i1}$$

Proof. The new transition matrix $\bar{\mathbf{P}}$ in its canonical form is:

$$\bar{\mathbf{P}} = \begin{bmatrix} (1-c)\mathbf{P} & \mathbf{R} \\ \mathbf{0}_{2 \times n} & \mathbf{I}_2 \end{bmatrix}, \quad \mathbf{R} \in \mathbb{R}^{n \times 2},$$

Where

$$\mathbf{R} = \begin{cases} c, & (i = u \wedge j = 1) \vee (i \neq u \wedge j = 2) \\ 0, & \text{otherwise} \end{cases}$$

We know[15] that absorption probabilities $\mathbf{B} = \mathbf{NR}$ where \mathbf{N} is the Fundamental matrix of process \bar{X} and as before $\mathbf{N} = \frac{1}{c}\mathbf{Q}$ So:

$$\begin{cases} \mathbf{B} = \mathbf{NR} \\ \mathbf{N} = \frac{1}{c}\mathbf{Q} \end{cases} \Rightarrow \mathbf{B}_{ij} = \sum_k \mathbf{Q}_{ik} \mathbf{R}'_{kj} = \begin{cases} Q_{iu}, & j = 1 \\ Q_i(V \setminus \{u\}), & j = 2 \end{cases}$$

□

Lemma 4.2. *For the personalized PageRank that the node i gives to itself, it holds:*

$$\mathbf{Q}_{ii} = c + (1-c) \frac{1}{k_i} \sum_{w \in E_i} (\mathbf{Q}_{wi})$$

Proof. From lemma 2.2 we know that $\mathbf{Q}_{iu} = \mathbf{B}_{i1}$. It also holds [15] that $\mathbf{B}_{i1} =$

$1\mathbf{R}_{ij} + \sum_{j \in E_i} \mathbf{P}_{ij}\mathbf{B}_{j1}$. So:

$$\begin{aligned}
\mathbf{Q}_{iu} &= \mathbf{B}_{i1} \\
&= \mathbf{R}_{ij} + \sum_{j \in E_i} (1-c)\mathbf{P}_{ij}\mathbf{B}_{j1} \\
&= \mathbf{R}_{i1} + (1-c) \sum_{j \in E_i} \mathbf{P}_{ij}\mathbf{Q}_{ju} \\
&= \mathbf{R}_{i1} + (1-c)\mu_{E_i}(\mathbf{Q}_{\cdot}(R)), \quad \mathbf{R}_{i1} = \begin{cases} c, & i = u \\ 0, & otherwise \end{cases}
\end{aligned}$$

□

From lemma 2.3 we can get now that :

$$k_u + 1 - \frac{(1-c)}{c}(\mathbf{Q}_{vu} - \frac{1}{k_u} \sum_{w \in E_u} (\mathbf{Q}_{wu})) = k_u - \frac{1}{c}[(1-c)\mathbf{Q}_{vu} - \mathbf{Q}_{uu}]$$

This quantity is always positive because $\mathbf{Q}_{vu} - \mathbf{Q}_{uu}$ is always negative and $1-c < 1$.

Lemma 4.3. $\forall v, u \in V$, it holds:

$$\mathbf{Q}_{vu} < \mathbf{Q}_{uu}$$

Proof. Let $f_{vu}^{(i)}$ be the possibility to reach transient state u starting from the transient state v for the first time at step i . And let $f_{vu}^* = \sum_{i=1}^{\infty} f_{vu}^{(i)}$. For the absorbing Markov chain \overline{X} it holds:

$$f_{vu}^* < 1 \tag{4.2}$$

That is because there is possibility c for transient state v to be absorbed at the first step to the absorbing state a_0 .

Let N_u to denote the number of visits to state u . then:

$$\begin{aligned}
P[N_u = m \mid \overline{X}_0 = u] &= f_{uu}^{*m-1}(1 - f_{uu}^*) \\
P[N_u = m \mid \overline{X}_0 = v] &= \begin{cases} 1 - f_{vu}^*, & m = 0 \\ f_{vu}^* f_{uu}^{*m-1}(1 - f_{uu}^*) \end{cases}
\end{aligned}$$

So N_u follows geometric distribution with success probability of $(1 - f_{uu}^*)$ and so:

$$\left. \begin{aligned} E[N_u \mid \bar{X}_0 = u] &= \frac{1}{1-f_{uu}^*} \\ E[N_u \mid \bar{X}_0 = v] &= f_{vu}^* E[N_u \mid \bar{X}_0 = u] \end{aligned} \right\} \Rightarrow$$

$$E[N_u \mid \bar{X}_0 = v] < E[N_u \mid \bar{X}_0 = u]$$

We also know that $E[N_u \mid \bar{X}_0 = v] = \mathbf{N}_{vu} = \frac{1}{c} \mathbf{Q}_{vu}$, $\frac{1}{c} > 0$ and so:

$$\mathbf{Q}_{vu} < \mathbf{Q}_{uu}$$

□

4.1.2 Impact of Adding a Set of Edges to a Single Source Node

The previous theorem can be generalized by adding a set of nodes to a fixed source node. In this case it holds the following.

Theorem 4.2. *If $G' = (V, E' := E \cup \tilde{E})$, $\tilde{E} = \{(u, v_i) \mid v_i \in V \wedge v \notin E_u \ i = 1, 2, \dots, \tilde{k}\}$, $E'_u = E_u \cup \tilde{E}_u$, $\tilde{E}_u = \{v \mid (u, v) \in \tilde{E}\}$ then:*

$$\mathbf{p}'(R) = \mathbf{p}(R) + \mathbf{p}_u \cdot \frac{\frac{(1-c)}{c} \left(\frac{1}{\tilde{k}} \sum_{v \in \tilde{E}_u} (\mathbf{Q}_v(R)) - \frac{1}{k_u} \sum_{w \in E_u} (\mathbf{Q}_w(R)) \right)}{\frac{k_u + \tilde{k}}{\tilde{k}} - \frac{(1-c)}{c} \left(\frac{1}{\tilde{k}} \sum_{v \in \tilde{E}_u} (\mathbf{Q}_{vu}) - \frac{1}{k_u} \sum_{w \in E_u} (\mathbf{Q}_{wu}) \right)}$$

4.1.3 Efficient Computation of Red Personalized PageRank and Personalized PageRank

From Lemma 2.2 we know that for the Markov chain \bar{X} , $\mathbf{Q}_{iu} = \mathbf{B}_{i1}$, $\mathbf{Q}_i(V \setminus \{u\}) = \mathbf{B}_{i2}$

Furthermore, we know that $\bar{\mathbf{P}}_{ij}^n$ gives as the probability to be in state j starting from state i after n steps. Also we know that $\lim_{n \rightarrow \infty} \bar{\mathbf{P}}^n$ exists. From the canonical form of $\bar{\mathbf{P}}$ we take:

$$\bar{\mathbf{P}}^n = \begin{bmatrix} (1-c)^n \mathbf{P}^n & \mathbf{R}^{(n)} \\ \mathbf{0}_{2 \times n} & \mathbf{I}_2 \end{bmatrix} \Rightarrow \mathbf{B} = \lim_{n \rightarrow \infty} \begin{bmatrix} \mathbf{R}^{(n)} \\ \mathbf{I}_2 \end{bmatrix} \Rightarrow$$

$$\mathbf{B} = \lim_{n \rightarrow \infty} \bar{\mathbf{P}}^n \cdot \mathbf{e}_{n+1}, \quad \mathbf{e}_{n+1} \in \mathbb{R}^{n+2} \quad (4.3)$$

The above expression allows us to compute the personalized PageRank of all nodes to node u with the computational cost of one PageRank.

Same as in Lemma 2.2 we can define the marcov chain \tilde{X} . We add two absorbing states $n+1, n+2$, from now on a_r, a_b respectively. We then unite all red nodes to state a_r with probability c and all the blue nodes to state a_b also with probability c . If $\mathbf{B}_{i1}, \mathbf{B}_{i2}$ are the absorbing probabilities for node i to a_r, a_b respectively and $\mathbf{Q}_i(R), \mathbf{Q}_i(B)$ is the ratio of personalized PageRank of node i for Red and Blue nodes, then we can prove the following.

Lemma 4.4. *The absorption probabilities of state i to state a_r, a_b of the absorbing Marcov chain \tilde{X} are equal to the ratio of personalized PageRank of node i that Red and Blue nodes receive respectively. This is:*

$$\mathbf{Q}_i(R) = \mathbf{B}_{i1}, \quad \mathbf{Q}_i(B) = \mathbf{B}_{i2}$$

Proof. proof is identical to lemma 2.2 □

Now working for \tilde{X} as before we can get an a similar expression for $\mathbf{Q}_i(R)$ and compute also the personilizes Red PageRank of all nodes.

From the above is clear that we need computational cost equal to three PageRankto compute the expression of theorem 2.1 for all possible future edges of the network with fixed source node.

4.1.4 Fair Important Edges in a Network

With the same methodology we can compute the formula for the deletion of an edge or a set of edges from a singe node:

Theorem 4.3. *If $G' = (V, E' = E \cup \{u, v\})$ then:*

$$\mathbf{P}(R)' = \mathbf{P}(R) - \mathbf{p}_u \cdot \frac{\frac{1-c}{c}(\mathbf{Q}_v(R) - \frac{1}{k_u-1} \sum_{w \in \tilde{E}_u} \mathbf{Q}_w(R))}{\frac{1}{k_u} - \frac{1-c}{c}(\frac{1}{k_u-1} \sum_{w \in \tilde{E}_u} \mathbf{Q}_{wu} - \mathbf{Q}_{vu})}$$

Theorem 4.4. *If $G' = (V, E' := E \cup \tilde{E})$, $\tilde{E} = \{(u, v_i) | v_i \in V \wedge v_i \notin E_u \ i = 1, 2, \dots, \tilde{k}\}$, $E'_u = E_u \cup \tilde{E}_u$, $\tilde{E}_u = \{v | (u, v) \in \tilde{E}\}$ then:*

CHAPTER 5

EXPERIMENTAL EVALUATION

5.1 Dataset Description

5.2 Fairness Aware PageRank Ranking

5.3 PageRank Fairness Aware Recommendations

5.1 Dataset Description

5.2 Fairness Aware PageRank Ranking

In this section, we evaluate experimentally the different fair Pagerank algorithms and provide quantitative and qualitative results. We have used various real data sets. We focus on the following four, while results for additional datasets can be found in the Appendix.

- **TWITTER:** A political retweet graph from [16].
- **DBLP2:** An author collaboration network constructed from DBLP including a subset of data mining and database conferences.
- **BOOKS:** A network of books about US politics where edges between books represented co-purchasing¹.
- **BLOGS:** A directed network of hyperlinks between weblogs on US politic [17].

¹<http://www-personal.umich.edu/~mejn/netdata/>

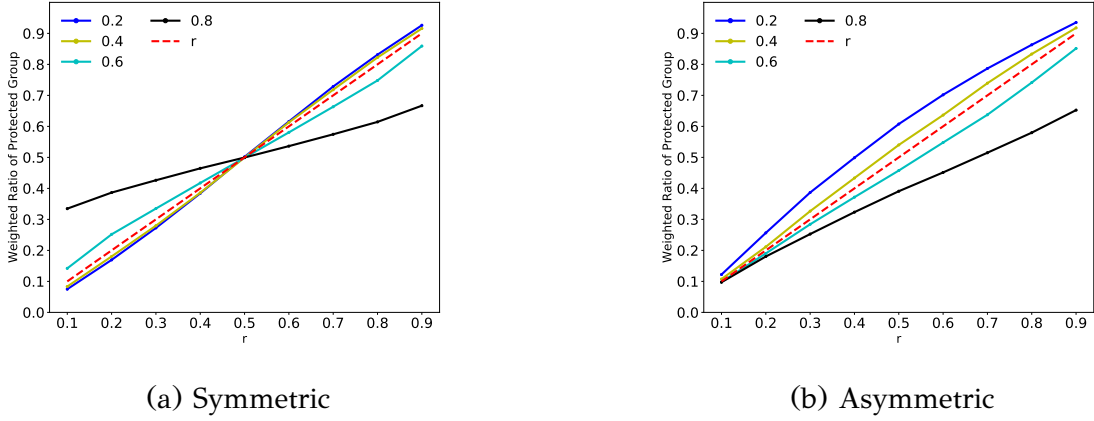


Figure 5.1: Fairness of the PageRank algorithm with the size of the protected group for varying homophily.

The characteristics of the real datasets, and the protected groups, are shown in Table 3.1. To infer the gender in the `DBLP2`, we used the python *gender guesser* package². We also report *homophily* which was shown to affect degree distributions among groups [8]. We measure it as the number of mixed edges, i.e., edges between nodes belonging to different groups, divided by $2r(r-1)$, i.e., the expected number of such edges. Values significantly smaller than 1 indicate that the network exhibits homophily [18].

Synthetic networks are generated using a variation of the biased preferential attachment model introduced in [8]. The graph evolves with time as follows. Let $G_t = (V_t, E_t)$ and $d_t(v)$ denote the graph and the degree of node v at time t , respectively. The process starts with an arbitrary initial connected graph G_0 , with $n_0 r$ red and $n_0(1-r)$ blue nodes. At each time step $t+1$, $t > 0$, a new node v enters the graph. The color of v is red with probability r and blue with probability $1-r$. Node v chooses to connect with an existing node u with probability $\frac{d_t(u)}{\sum_{w \in G_t} d_t(w)}$. If the color of the chosen node u is the same with the color of the new node v , then an edge between them is inserted with probability h ; otherwise an edge is inserted with probability $1-h$. If no edge is inserted, the process of selecting a neighbor for node v is repeated until an edge is created.

Parameter h controls the level of homophily in the network, where $h = 0$ corresponds to homophily, $h = 0.5$ to the random case and $h = 1$ to heterophily. We also consider asymmetry in homophily. In this case, the above procedure is followed by

²<https://pypi.org/project/gender-guesser/>

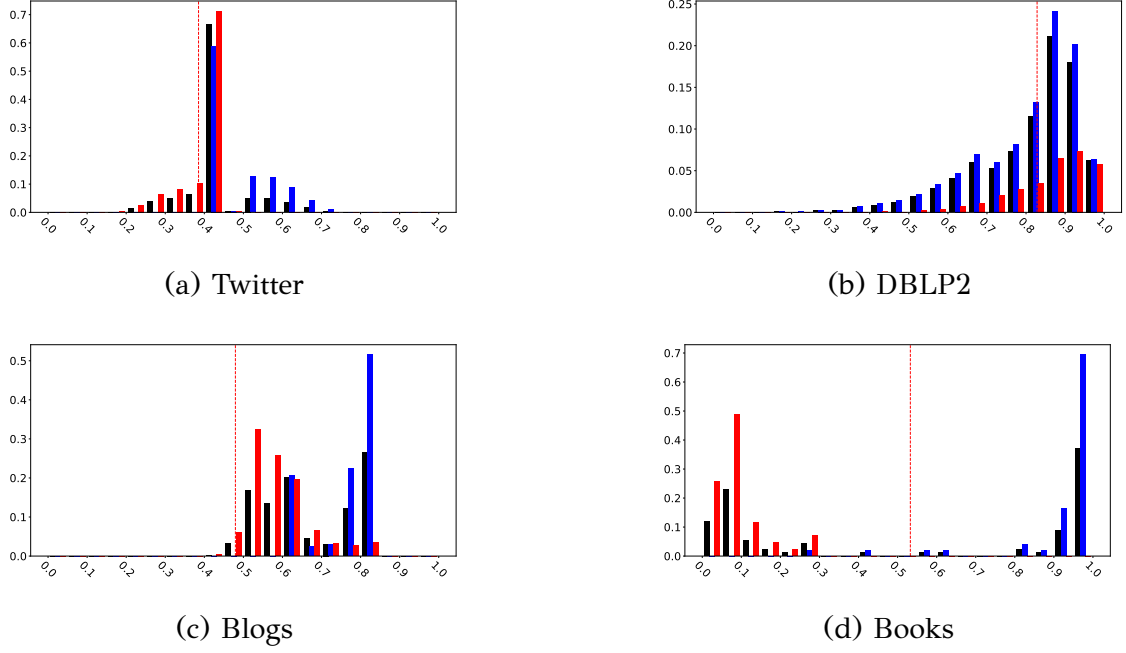


Figure 5.2: Personalized pageranks starting from each of the blue nodes: histogram of the fraction of the personalized weight. The x-axis corresponds to weights fractions (blue, red and all (black bar)) and the y-axis to the percentage of the blues nodes with the corresponding fraction. The majority of nodes allocate the larger fraction of the personalized weight to blue nodes, thus being highly unfair to the opposite group.

a node v only when v belongs to the red group. A node v in the blue group connects with the selected node u without testing u 's color.

The datasets and code are available at GitHub ³.

Fairness in the original Pagerank algorithm. We use the synthetic datasets to study the behavior of Pagerank for different levels of homophily and relative sizes of the two groups. For this set of experiments, we set $\phi = r$. As shown in Figure 5.1, for the symmetric case, when the groups exhibit homophily ($h = 0.2$ and $h = 0.4$), PageRank is unfair towards the minority group. On the contrary, when the groups exhibit heterophily ($h = 0.6$ and $h = 0.8$), then PageRank is unfair towards the majority group. For the asymmetric case, i.e., when the blue group shows no homophily, being homophilic helps the red group independently of its size, while being heterophilic hurts the red group independently of its size.

³<https://github.com/SotirisTsioutsoulis/FairLaR>

Table 5.1: Utility loss with respect to optimal utility ($\frac{LFPR_X}{OPTIMAL}$, for $\phi = 0.5$)

Dataset	LFPR _N	LFPR _U	LFPR _P	LFPR _O	SFPR
TWITTER	6.576	6.683	4.218	2.1671	2.699
DBLP2	1.356	1.232	1.516	1.1792	2.6
BLOGS	5.05	5.08	3.163	1.5923	1.73
BOOKS	9.53	4.94	1.576	1.000	1

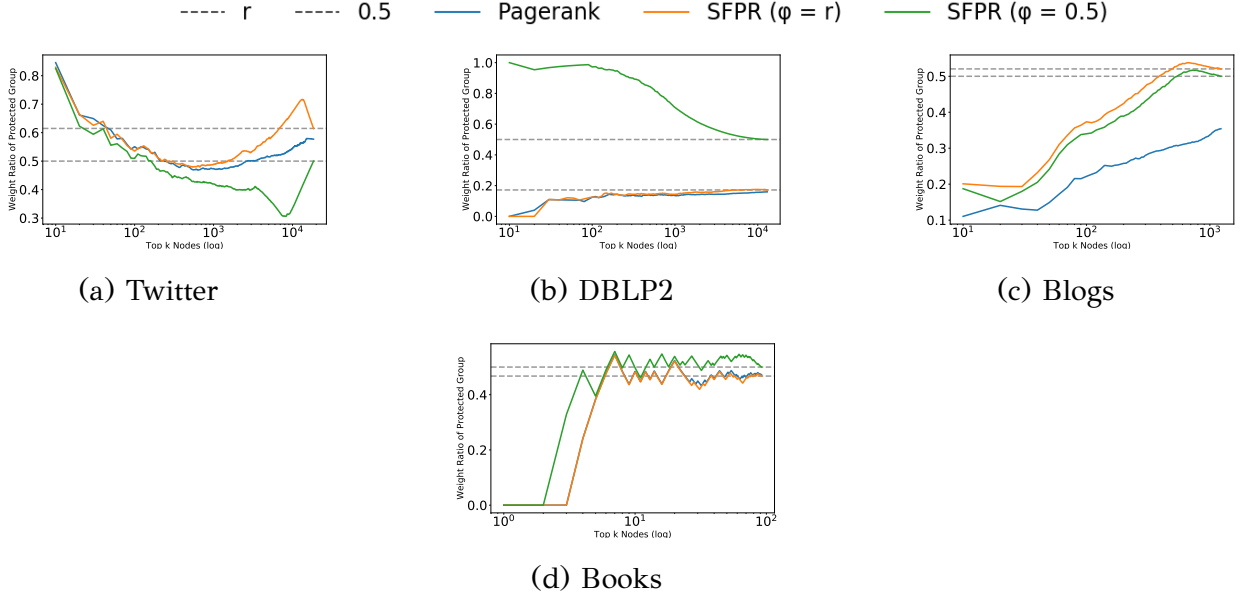


Figure 5.3: Fairness sensitive Pagerank for $\phi = r$ and $\phi = 0.5$.

For the real dataset, we report the fraction of the total weight allocated to each of the two groups in Table 3.1. In some cases (BLOGS, TWITTER), the fraction of the weight assigned to the protected group is significantly smaller than r . In all cases, by setting ϕ to the desired level of fairness, we can redistribute weights so that we get the desired ϕ -fairness. We report quantitative and qualitative results for $\phi = 0.5$ in the next section.

To get a better insight about the distribution of the weights between the two groups, we also run personalized Pagerank algorithms starting from each node i and calculated for each node i the fraction of the weight allocated to the blue and red nodes (ignoring the Pagerank allocated to the node itself). In all graphs, most of the starting nodes allocate the majority of their personalized pagerank weights to nodes in their group, resulting in highly unfair weights. We report the histogram of the fraction of the weights allocated to blue and red node for personalized pageranks

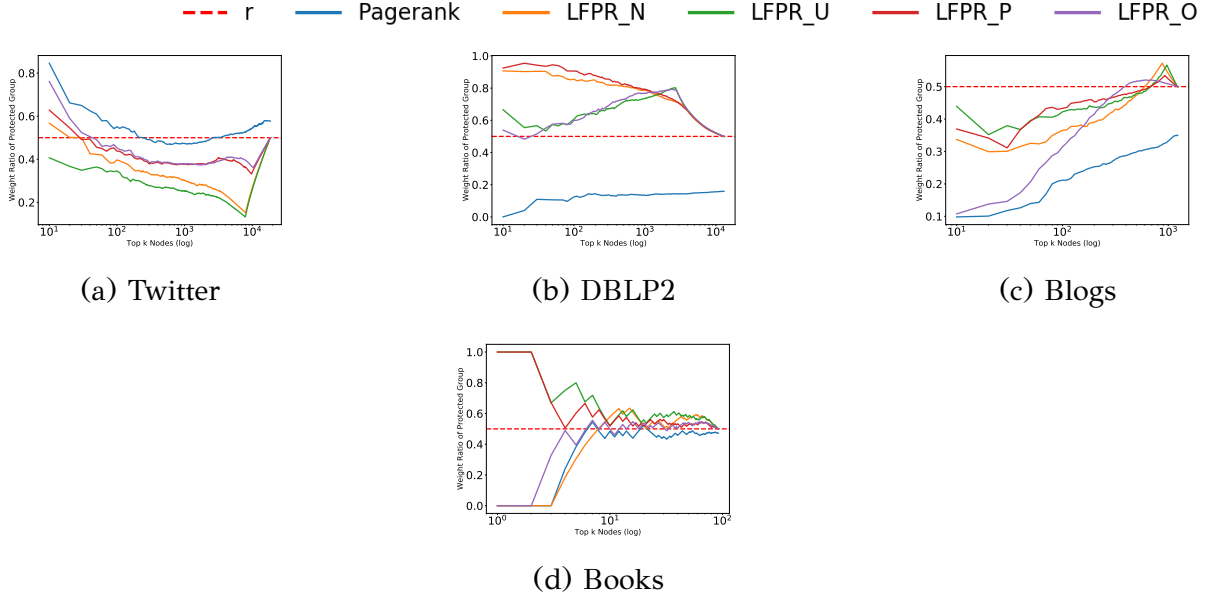


Figure 5.4: Locally fair Pagerank algorithms for $\phi = 0.5$.

starting from each of the blue nodes in Figure 5.2. Correlation with homophily can be observed, with the most homophilic networks, i.e, BOOKS and TWITTER, showing the largest unfairness. Our locally fair Pagerank algorithms can be used to attain ϕ -fairness for personalized Pagerank as well.

The fair PageRank algorithms. We run our fair Pagerank algorithm for various values of ϕ . In Figure 5.3, we report results for $\phi = r$ and $\phi = 0.5$ for the fairness sensitive Pagerank (SFPR), while in Figure 5.4, results for $\phi = 0.5$ for the various locally fair Pagerank algorithms (i.e, neighbor (LFPR_N), uniform (LFPR_U), proportional (LFPR_P) and with optimized residual (LFPR_O)). Results for $\phi = r$ can be found in the Appendix.

Table B.2 reports the utility loss for each of the fair pagerank algorithms relative to the optimal utility loss as estimated by Algorithm 3.1. For the non-optimized algorithms, as expected taking into account the original pagerank values, the LFPR_P algorithm results in the smallest utility loss. The LFPR_N algorithm incurs the highest utility loss. The utility loss decreases significantly when considering the optimized algorithms. It is interesting that for different datasets different variants perform better. This suggests that the different algorithms provide different levers for adjusting fairness. Depending on the dataset one approach may be more applicable for preserving utility than another. We discuss this further in our qualitative comparison of the algorithms.




Figure 5.5: Targeted fair PageRank algorithms and the optimal post-processing redistribution for $\phi = 0.5$. In each dataset, the target set S includes the 10% of nodes with the highest PageRank weights.

We report the results of all targeted fair PageRank algorithms in Figure 5.5. The targeted fair PageRank algorithms allows us to focus on a specific set of nodes and adjust their weights in a fair manner. In this experiment, we selected the set S for each dataset, so as to include the 10% of the nodes having the highest original PageRank values. This provides us with the flexibility to adjust weights in the top positions.

Residual distribution policies. In this set of experiments, we take a closer look at the set of nodes that are mostly affected by the different residual distribution policies. To this end, we consider the sets LOSS (resp. GAIN) with the 10 red and 10 blue nodes whose PageRank decreased (resp. increased) the most. For each node i , we define its *in-neighborhood fairness*, $in_f(i)$ as the ratio of its red in-neighbors over all its in-neighbors. Thus, $in_f(i) = 0$ corresponds to a node i with only blue in-neighbors, $in_f(i) = 1$ to only red in-neighbors and $in_f(i) = 0.5$ to a balanced in-neighborhood.

We run the algorithms on all datasets (including the ones in the Appendix). The results are depicted in Figure 5.6. As shown in Figure 5.6(a), with very low variance, all algorithms penalize the nodes that are in homophilic in-neighborhoods, that is, red




Figure 5.6: In neighborhood fairness for the nodes with the maximum loss and gain for each algorithm.

nodes with large in_f , and blue nodes with small in_f . The algorithms exhibit different behavior regarding which nodes each algorithm promotes as shown in Figure 5.6(b). $LFPR_N$ promotes the red (i.e., protected group) nodes that are in heterophilic (i.e., protected-group dominated) in-neighborhoods. This does not hold for the blue nodes. On the other hand, $LFPR_P$ does not seem to particularly favor heterophilic red nodes, while it tends to follow the trend of the PageRank promoting homophilic blue (i.e., favored-group) nodes. Finally, for $LFPR_U$, in_f seems to play a lesser role, with the small favoritism to homophilic blue nodes most probably reflecting the homophily of the blue group in the underlying population.

Qualitative comparison. To provide some insight on the weights produced by the various algorithm, we visualize their output for $\phi = 0.5$ in Figures 5.7 and 5.8. In the visualizations, red nodes are colored red, and blue nodes are colored blue. Their size depends on the value of the quantity we visualize.

For the TWITTER and the BOOKS datasets, where the fraction of the weight of the protected group is close to ϕ , the fairness sensitive pagerank is very similar to the original one. For the BLOGS and especially for the DBLP2 datasets, where the fraction of the weight of the protected (red) group is much smaller, the fairness sensitive pagerank promotes red nodes. We also visualize the jump vector for the fairness sensitive pagerank. We observe that for the TWITTER and the BOOKS dataset, where the algorithm is already “almost” fair, the jump vector assigns rather uniform weights, as the original Pagerank. For the other two datasets, it gives large values to a number of red nodes. This suggests an interesting line for future work: considering these nodes in link recommendation algorithms, since it seems that these nodes play a role in

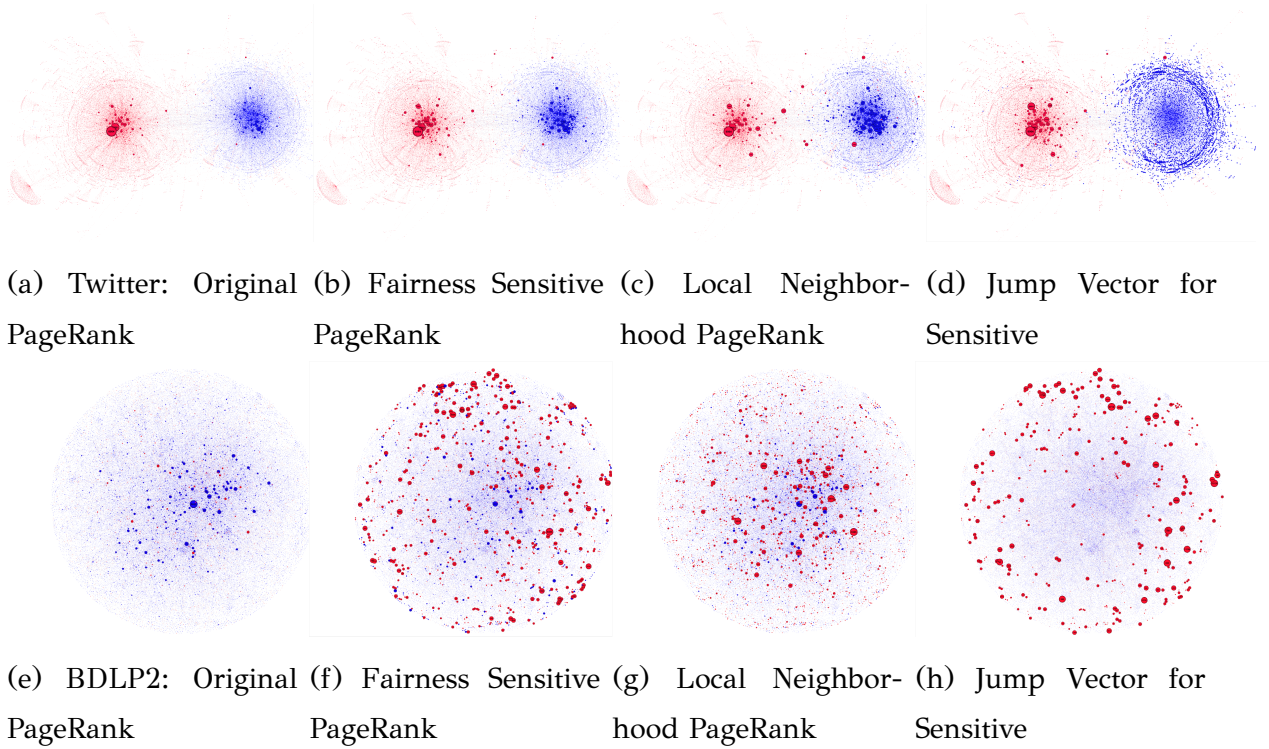


Figure 5.7: Visualization for $\phi = 0.5$

fairness.

The neighborhood locally fair pagerank algorithm produces different weights from the original Pagerank for all four datasets. In all cases, it promotes nodes connecting the two opposite groups, i.e., nodes that are minorities in their neighborhoods. This is more evident in the most homophilic networks, that is, in `TWITTER` and `BOOKS`. Such nodes are also known as weak links and play an important role. They can also be useful in the context of recommendations, since research shows that it is more likely for such nodes to be accepted from the other side [19].

5.3 PageRank Fairness Aware Recommendations

In this section we compare and evaluate different known recommendation policies and our novel ones. We start by studying various known link recommendation systems and we show that they don't improve our objective (enhance the protected group in the network). Due to the fact that our policy acts complementary to an existing one we could use any recommendation policy to apply it on. In these experiments we use `node2vec` recommender as it is one well known recommender that has been used

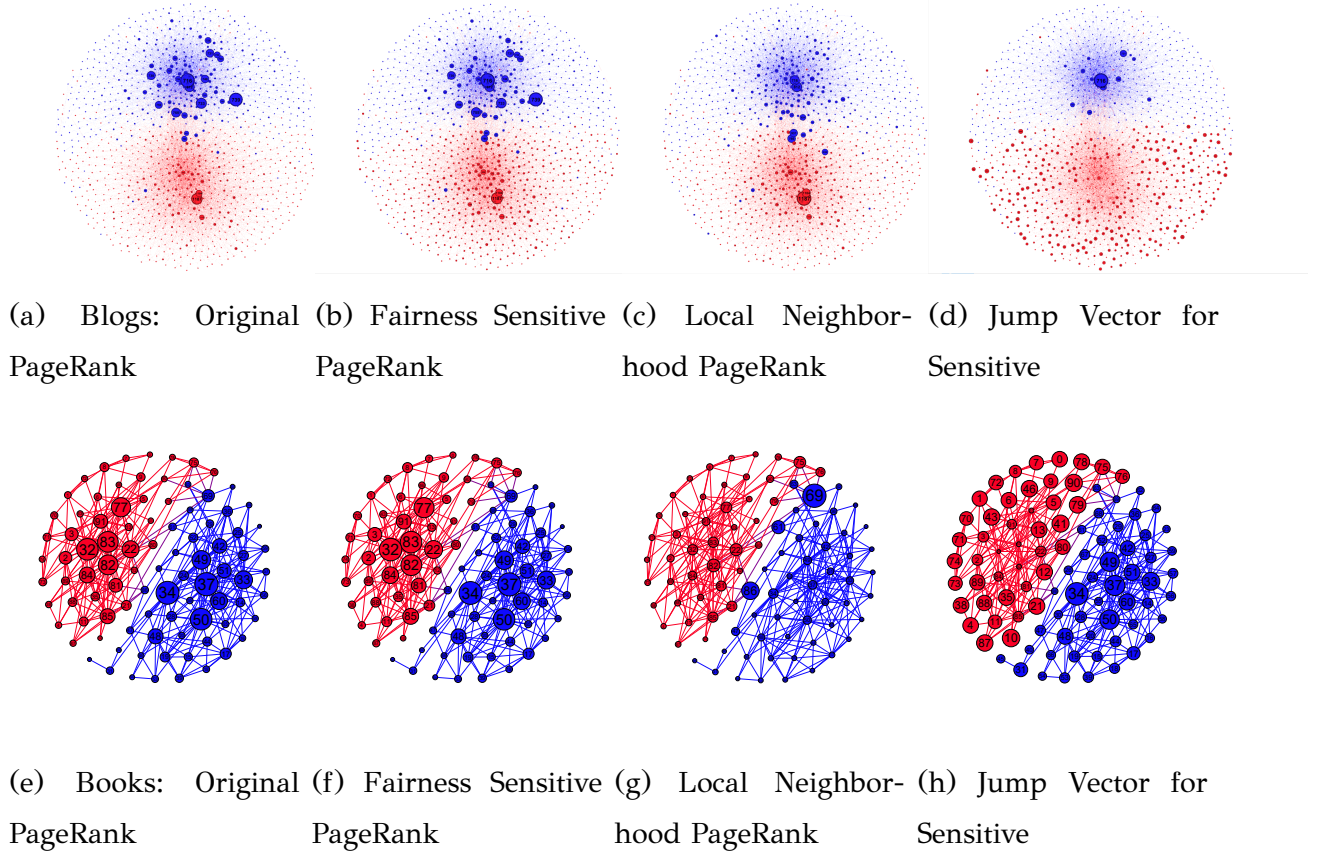


Figure 5.8: Visualization for $\phi = 0.5$

extensively in research with good results.

Node2vec recommender was implemented by taking the node2vec implementation of snap ⁴ and train a logistic regression classifier with sklearn module in python. For both the classifier and the node2vec embeddings we used the default settings. For the training of the classifier we use as train test the 80% of the network's edges for positive examples and equal amount of edges that don't exist for negative example. We use the rest 20% of positive edges and equal amount of negative edges as test set. For all the rest recommendation algorithms we used their implementation on networkx ⁵ module for python

To evaluate our policy we compare both the impact to our objective and the quality of link recommendations it produces. We evaluate the quality of recommendations by assuming that the acceptance probability of candidate links coming from the original recommender (node2vec in our case) is in fact true. We proceed further the analysis of our policy by identifying the differences in quality features of link recommendations

⁴<http://snap.stanford.edu/snap/index.html>

⁵<https://networkx.github.io/>

and explaining in a level the underlying mechanism of each one.

5.3.1 Classic Recommendation Policies

Each of the link recommendation policies has its own mechanism of selecting and proposing links to a source node. Also, as we saw in theorem 4.1, that the impact of a link addition to the network it depends not only on the node that is being proposed but also on the node that is proposed to. To study the impact of the different policies to the red PageRank of the network we have decided to use the resource allocation, Jaccard coefficient, Adamic Adar, preferential attachment and node2vec recommendation policies. In this direction we conducted the following experiment.

First we choose a set of source nodes. Then we take the 10 best link recommendations for every source node by each policy. After that, we add one edge per source node (starting from the best one and continuing with the second best etc.). We continue the process until we have added all the 10 edges to all source nodes. We use K to denote the number of the links that have been added to each of the source nodes (e.g. $K = 4$ means that we have added 4 links to each of the source nodes). We conduct the experiments for 3 different sets of source nodes. In the first one we chose randomly 10% of the nodes and for the other two we chose a hundred best red and a hundred best blue nodes. By best we mean those nodes in which we expect to have the greatest expected impact on the network as derived from the formula in theorem ???. The expected impact for the source nodes is computed by approximation due to the prohibitive complexity of the actual computation. This approximation is $(PageRank)/(outdegree)$. We present the red PageRank of the network as evolves after the addition of the recommended edges for $K = 0, 1, \dots, 10$.

In general, we can separate the known recommendation policies into two categories. In the first we classify policies that their recommendation score is based on the number of the common neighbors between two nodes (resource allocation index, Jaccard coefficient, Addamic Adar) and In the second one policies that tend to propose strong/popular nodes in the network (preferential attachment, node2vec). Algorithms in each category tend to share common characteristics and behavior, or equivalently they affect the networks evolution in the same way.

To start with, we can see in fig. 5.9 that policies of the first group behave similar with the random policy in a smooth way, giving a small privilege to the bigger group.

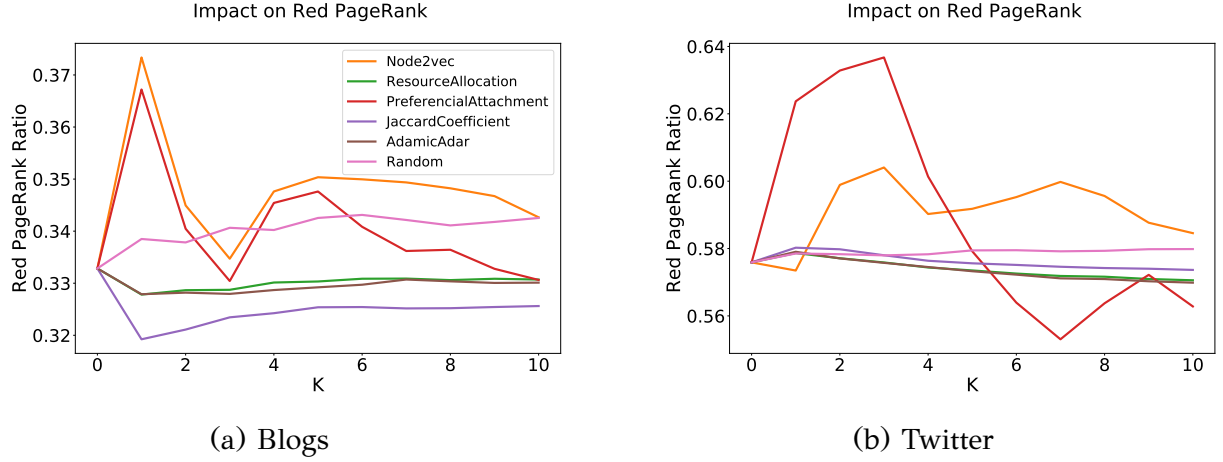



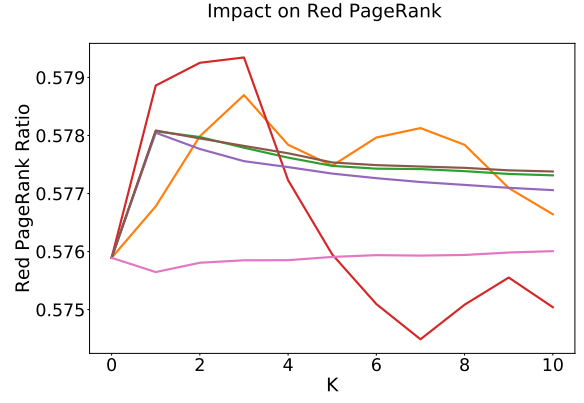
Figure 5.9: Impact on Fairness - Known Recommenders - Random Source Nodes.

The policies of the second group have a slightly unpredicted behavior in the first 1 - 3 steps but they too follow the trend of the initial PageRank as the network evolves. This unpredicted behavior can be explained - as we will see below - due to the fact that policies of the second group, in contradiction with the ones of the first (see table 5.2), tend to propose common neighbors independently of the source node. Thus, the characteristic of the first few proposed nodes dominate the impact on the network. As the number of the links accepted added to the source node gets higher, the more the initial trend of the network is represented in the target nodes, so the impact of the proposed links in networks Red PageRank approaches that of the random link recommendations.

As far as the source node concerns, we can see from figures 5.10, 5.11 that the behavior of the policies of second group change slightly in the pace they converge to the random policy, meaning they are more robust to source node selection while the policies of the first group change total behavior and they enhance the team that the source nodes belong to. This is expected, if we consider that social networks tends to be homophilic (this also holds in our datasets) by their nature and these kind of recommendation policies propose links in distance two. The first one means that a blue source node will probably have blue neighbors and it will exists in a blue dominated cluster and that a blue node will also have a higher personalized blue PageRank. The equivalent holds for a red source node. So from our formula about edge addition impact and considering the second observation about the distance of proposed links, it is clear that the links proposed of the policies of first group are ideal to rise the PageRank of the source node's group.




(a) Blogs




(b) Twitter

Figure 5.10: Impact on Fairness - Known Recommenders - Red Source Nodes.



(a) Blogs



(b) Twitter

Figure 5.11: Impact on Fairness - Known Recommenders - Blue Source Nodes.

To sum up, though the link recommendation policies already exists have shown some valuable results, proposing edges in a simple way and having an expected impact on the network, they can not be used as link recommendation system if we care to restrict the discrimination on a network, improve fairness or in more general term enhance the presence of a protected group. We continue the experimental evaluation by showing how we can tackle this behavior and how our proposed fair policy can affect a recommender towards this direction.

5.3.2 Fair Recommendation Policies

In this section we study the link recommendations provided by our fair recommender, from node2vec recommender and from the hybrid fair recommender. We observe




Figure 5.12: Impact on Fairness - Fair Recommenders - Random Source Nodes.

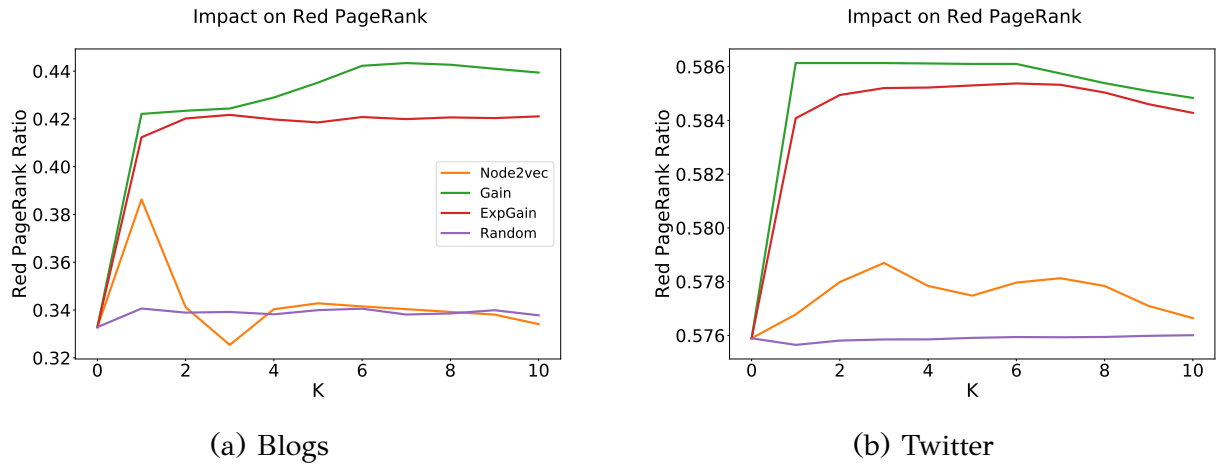


Figure 5.13: Impact on Fairness - Fair Recommenders - Red Source Nodes.

that between node2vec and the fair recommender there is a significant trade off between the quality and the wanted impact of the links, however this trade off is nicely balanced by the hybrid fair recommender.

To compare this result we present the red PageRankratio of each network as we did before, only this time for random, node2vec, fair and hybrid fair recommenders and for the same experiment we also present the average acceptance probability of the networks as this calculated based on node2vec recommendation score.

As we see in figures 5.12, 5.13, 5.14, the fair policy rises the red PageRank more than any other and the hybrid fair policy succeeds impressive rise as well. This shows us that taking into consideration both fair and recommendation score doesn't affect significant the wanted impact on the network. Moreover by figures 5.15 5.16 5.17 we conclude that fair policy doesn't take into consideration at all the acceptance

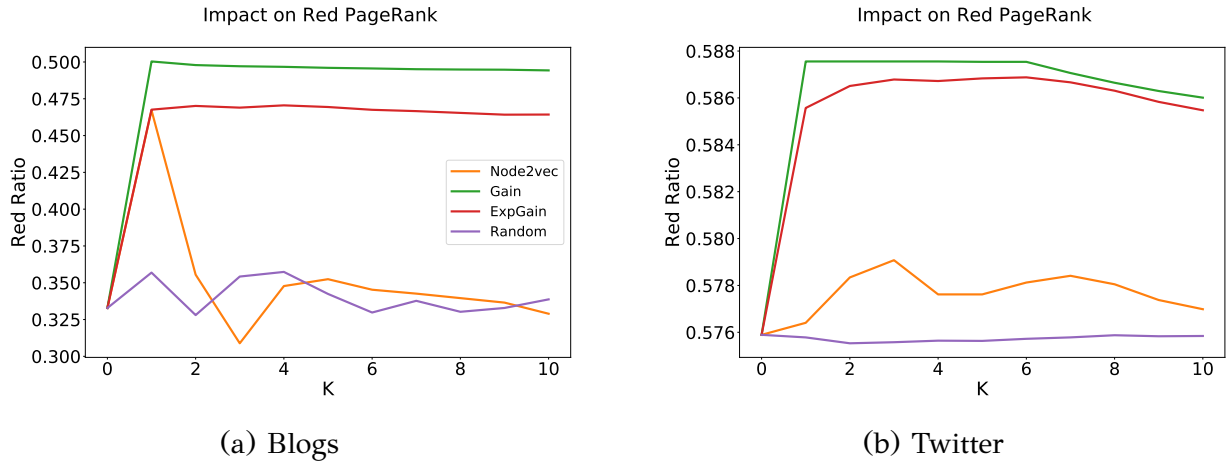


Figure 5.14: Impact on Fairness - Fair Recommenders - Blue Source Nodes.


probability of the link it proposes, performing in some cases worst than the random recommendation policy. This makes these recommendation invaluable as it is highly possible not be accepted by the users. On the other side we see that the hybrid fair policy manages to restrict this phenomenon, approaching in a satisfying level the average recommendation score of the node2vec recommender. Last but not least we see that the results are not affected by the set of source nodes which is really important as in practice social networks differs in their demographic characteristics.

An interesting observation about the impact of fair recommendation policies is presented in figures 5.18 5.19. In this figures we see the ratio of red PageRank at top k nodes by PageRank as it has been formed at the end of the previous experiment. The interesting effect of fair policies in contrast with the node2vec is that fair policies improve the red PageRank ratio in a network while node2vec preserve the original distribution. This difference means that fair recommendations help the protected group to gain higher scores and positions in the ranking by PageRank algorithm and so be represented more fairly in the top positions.

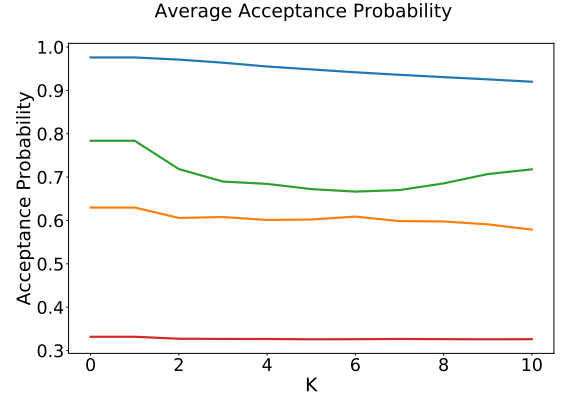
Having this in mind, it seems like the hybrid fair policy seems a good solution out our problem. We continue by studying the

5.3.3 Target Nodes Analysis

So far it seems that hybrid fair policy performs really good in all the metrics of our evaluation. However, having a link recommender in use, is always interesting to understand its link recommendation mechanism in a more simple way. That is

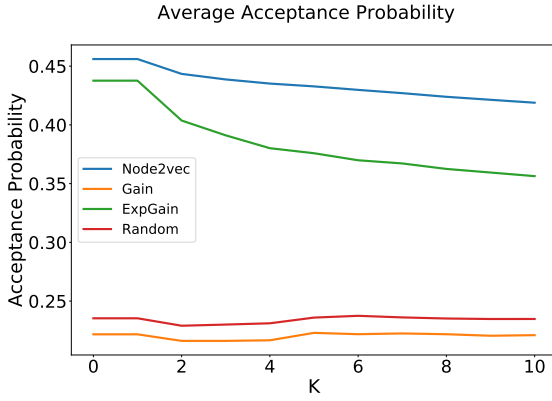


(a) Blogs




(b) Twitter

Figure 5.15: Average Acceptance Probability - Random Source Nodes.



(a) Blogs



(b) Twitter

Figure 5.16: Average Acceptance Probability - Red Source Nodes.

understanding what are these node and edge characteristics that rule its decisions and which is the dynamic of each proposed target node to the network. This kind of results most of times aren't so unexpected and they provide us with useful insights based on simple metrics. In this direction we study the nodes that have been selected from the previous experiment. We first present in table 5.2 the total distinct number of nodes for each policy per dataset and then, in tables 5.3, 5.4 we present the quality features for the node2vec, fair and hybrid fair policies.

In table 5.2 we see that node2vec, preferential attachment, fair and hybrid fair,

Table 5.2: Number of Total Unique Target Nodes by Policy

Dataset	Node2vec	Fair	Hybrid Fair	Resource Allocation	Jaccard Coefficient	Adamic Adar	Preferential Attachment
Blogs	15	20	40	351	662	315	17
Twitter	51	154	105	5184	7492	4833	12

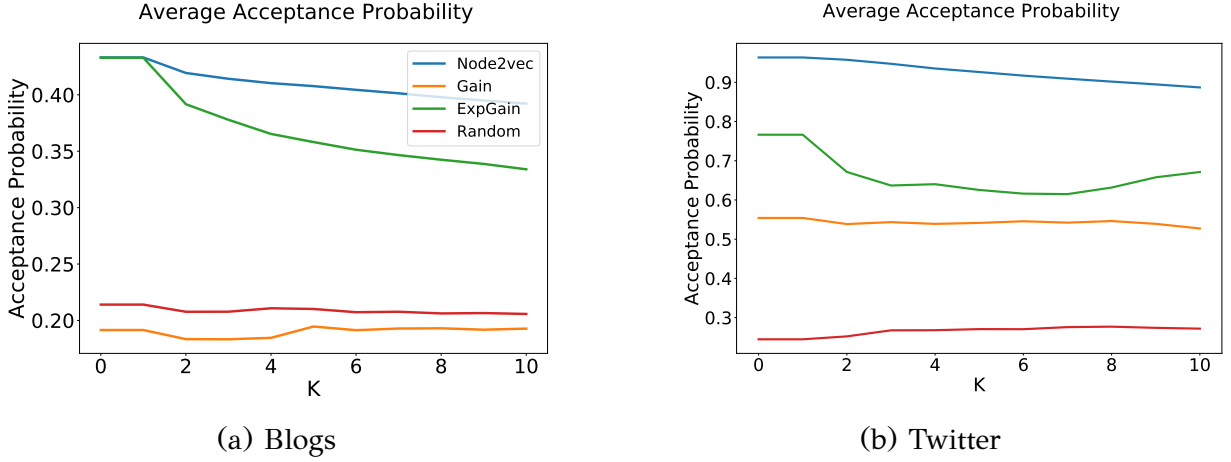


Figure 5.17: Average Acceptance Probability - Blue Source Nodes.




Figure 5.18: Red PageRank ratio at top k nodes by PageRank - Blogs network.

are tend to propose a relative small number of distinct target nodes. On the other, side the rest of the policies propose a much larger number. This result follows our instinct as the node2vec, preferential attachment and fair policies tends to highlight nodes strong globally in the network while the rest policies propose nodes that are locally strong. The result for the hybrid fair policy can be explained straight forward as it is the combination of the node2vec and fair policies.

Except the fact that node2vec, fair and hybrid fair policies select a relative small number of distinct target nodes to propose, there is a subset of them that have significant more occurrences. To study further the 3 policies we keep only the top in occurrences nodes. The selection process is described in figures 5.20, 5.21. The exact number of minimum occurrences that we accept is found by plotting the occurrences of each node in descending order, we observe that this plot approaches a sigmoid function, something that allow us to define properly this minimum.

We present these results in tables 5.3, 5.4. First we observe that node2vec rec-

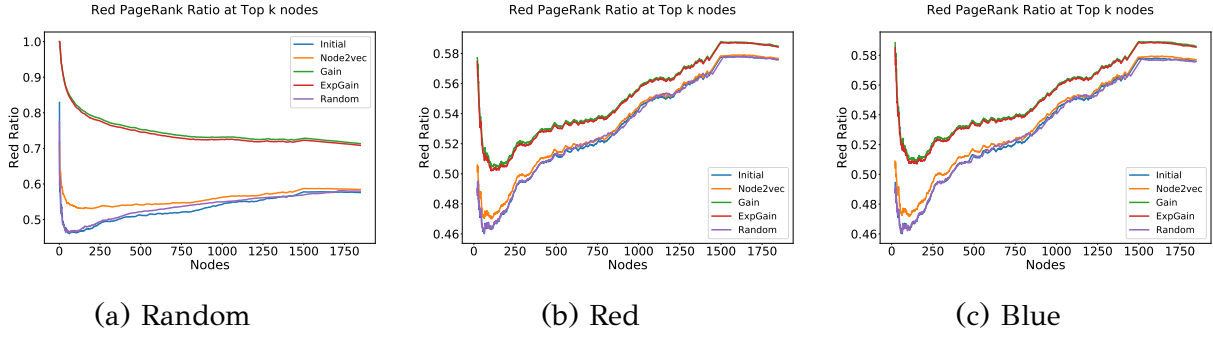


Figure 5.19: Red PageRank ratio at top k by nodes PageRank - Twitter network.

Table 5.3: Target Quality Features in Blogs Network.

Policy	Distance			PageRank			Red PageRank			Node Homophily		
	mean	median	max	mean	median	max	mean	median	max	mean	median	max
Random	2.809016	3	7	0.000243	0.000331	0.045172	0.000000	0.282878	0.638946	0.000000	0.500000	1.000000
Node2vec	2.313158	2	4	0.000243	0.004722	0.010006	0.161032	0.313908	0.564971	0.000000	0.099480	0.957143
Fair	3.745805	4	7	0.000243	0.000243	0.000583	0.615104	0.620985	0.638946	1.0	1.0	1.0
Hybrid Fair	2.644818	3	5	0.000243	0.001271	0.006028	0.490224	0.555863	0.638946	0.834951	0.970612	1.000000

ommendations are characterized by nodes that are strong by PageRank and gather the distribution of distances around smaller values. Fair policy nodes are high in red personalized PageRank and in red out neighbors ratio while hybrid fair balances all the above having in general higher values in all the forthmentioned scores.

5.3.4 Homophily and Minority Size

We have already seen in figure ?? the effect of homophily and minority size to the Red PageRank ratio of a network. A subsequent question is if and how these factors affect the evolution of Red PageRank in a network depending on the the link recommendation policy.

We present in figures 5.22, ?? the impact of the different recommendation policies in the evolution of the networks for different degrees of protected group size and different degrees of symmetric and asymmetric same group preference probability. Synthetic network confirm the behavior we show in the subsection 5.3.2 that node2vec

Table 5.4: Target Quality Features in Twitter Network.

Policy	Distance			PageRank			Red PageRank			Node Homophily		
	mean	median	max	mean	median	max	mean	median	max	mean	median	max
Random	4.98	5	12	0.000054	0.000037	0.003275	0.577001	0.639552	1.0	0.512503	0.5	1.0
Node2vec	3.85	4	11	0.001410	0.001376	0.003275	0.643997	0.733387	0.817765	0.750000	1.0	1.0
Fair	5.27	5	11	0.000185	0.000228	0.000298	0.942876	1.0	1.0	1.0	1.0	1.0
Hybrid Fair	4.87	5	11	0.000457	0.000283	0.001412	0.935682	1.0	1.0	1.0	1.0	1.0

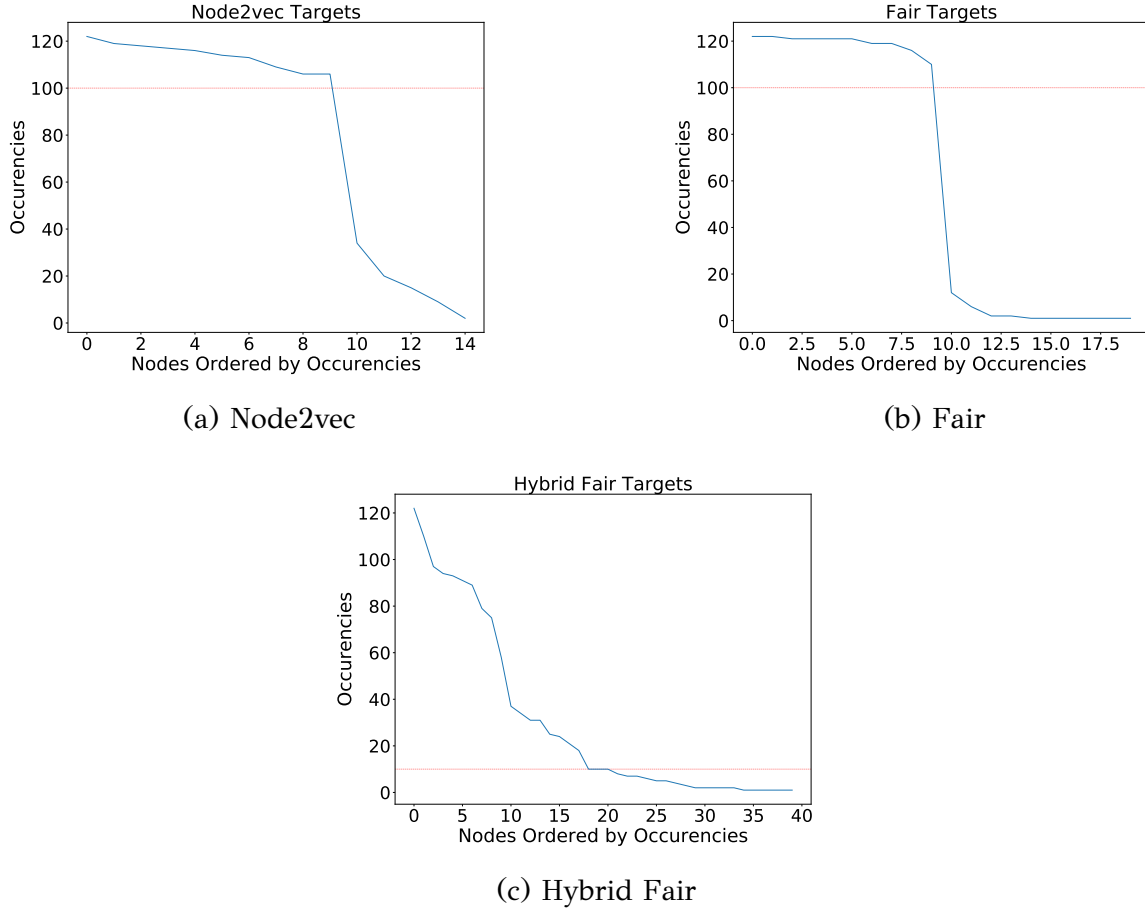


Figure 5.20: Cutting Point for Selecting Nodes.

doesn't change the network's red PageRankratio. Fair policies they both enhance protected group as expected but we see that their dynamics are affected both by the size and the same group preference probability. The general rule that applies to both forth mentioned quantities is the greater the values the greater the impact.

5.3.5 Batch vs Online Gain

Fair policy exhibits satisfactory results but we can understand from the formula that it doesn't take into consideration the changes in the network's structure that happens from the addition of edges on other nodes. We conduct the basic experiment with a greedy algorithm recalculating the fair score before any recommendation. Figures 5.23 5.24

We see that in many cases the greedy algorithm can extend in grade level the impact on the network but there are also cases where the two algorithms perform

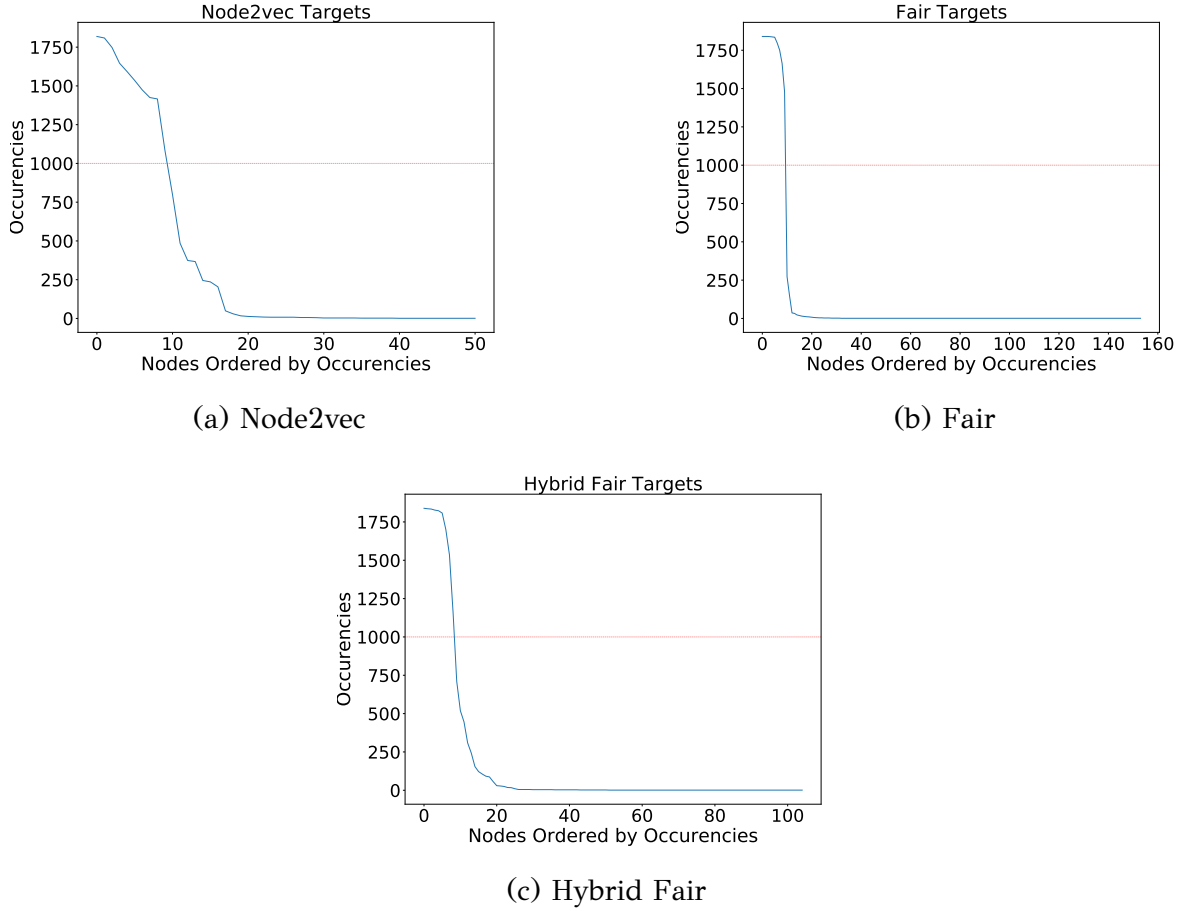


Figure 5.21: Cutting Point for Selecting Nodes

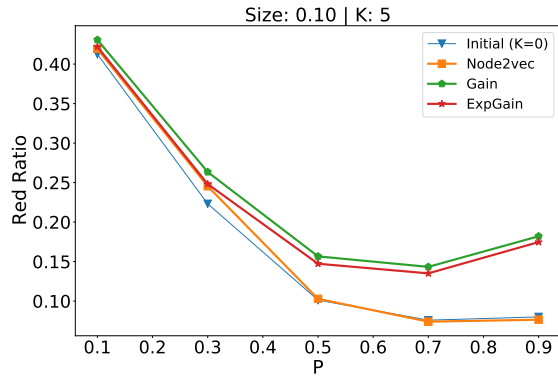
equivalently. This happens when the changes in the network doesn't change the order of best targets per source.

5.3.6 Fairness, Accepted Probability Correlation


From the experiments so far it is obvious that it is difficult to combine high recommendation score with high fair impact. In fact, the next experiments show us that there is also a negative correlation between acceptance probability and fair score on the top suggestions of each.

To measure that we use the recommendations of node2vec and fair policy we had for the synthetic networks. From these sets we keep the best 50%. For the edges proposed from node2vec recommender we separate them in buckets of equal size and we plot the average fair score for all buckets. We create the corresponding plots for the edges proposed from fair recommender. The number and the size of buckets


differs in every occasion depending on the range of the values. Results are presented in figure 5.25




(a) Node2vec




(b) Node2vec



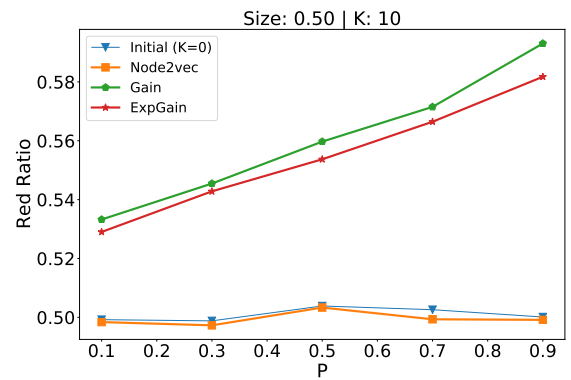
(c) Node2vec



(d) Node2vec



(e) Node2vec



(f) Node2vec

Figure 5.22: Red PageRank ratio to different same group preference probability for sizes 0.1, 0.3, 0.5 after 5 and 10 link additions for random sources.




Figure 5.23: Fairness Impact for Batch and Online Fair Policy - Blogs.




Figure 5.24: Fairness Impact for Batch and Online Fair Policy - Twitter.

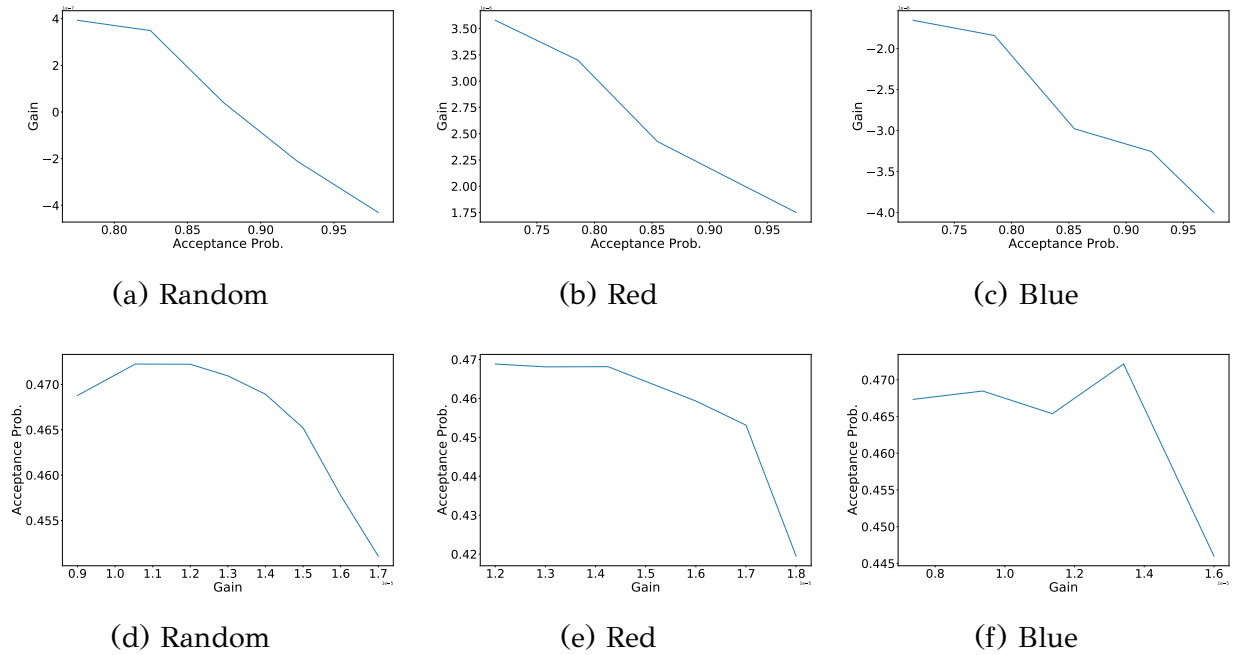


Figure 5.25: Recommendation Score - Fair Score Correlation.

CHAPTER 6

RELATED WORK

Algorithmic fairness. Recently, there has been increasing interest in algorithmic fairness, especially in the context of machine learning. Fairness is regarded as the lack of discrimination on the basis of some protective attribute. Various definition of fairness having proposed especially for classification [13, 1, 20, 21]. We use a group-fairness definition, based on parity. Approaches to handling fairness can be classified as *pre-processing*, that modify the input data, *in-processing*, that modify the algorithm and *post-processing* ones, that modify the output. We are mostly interested in in-processing techniques.

There is also prior work on fairness in ranking [22, 23, 24, 25]. All of these works consider ranking as an ordered list of items, and use different rules for defining and enforcing fairness that consider different prefixes of the ranking [22, 23], pair-wise orderings [25], or exposure and presentation bias [26, 24].

Our goal in this paper is not to propose a new definition of ranking fairness, but rather to initiate a study of fairness in link analysis. A distinguishing aspect of our approach is that we take into account the actual Pagerank weights of the nodes, not just their ranking. Furthermore, our focus in this paper, is to design in-processing algorithms that incorporate fairness in the inner working of the Pagerank algorithm. We present a post-processing approach as a means to estimate a lower bound on the utility loss. None of the previous approaches considers ranking in networks, so the proposed approaches are novel.

Fairness in networks. There has been some recent work on network fairness in the context of graph embeddings [27, 28, 29]. The work in [27] follows an in-processing

approach that extends the learning function with regulatory fairness enforcing terms, while the work in [28] follows a post-processing approach so as to promote link recommendations to nodes belonging to specific groups. Both works are not related to our approach. The work in [29] extends the node2vec graph embedding method by modifying the random walks used in node2vec with fair walks, where nodes are partitioned into groups and each group is given the same probability of being selected when a node makes a transition. The random walk introduced in [29] has some similarity with the random walk interpretation of LFPR_N. It would be interesting to see, whether our extended residual-based algorithms could be utilized also in the context of graph embeddings, besides its use in link analysis.

There are also previous studies on the effect of homophily, preferential attachment and differences in group sizes. It was shown that the combination of these three factors leads to uneven degree distributions between groups [8]. Recent work shows that this phenomenon is exaggerated by many link recommendation algorithms [9]. Evidence of inequality between degree distribution of minorities and majorities was also found in many real networks [30]. Our work extends this line of research by looking at Pagerank values instead of degrees.

Finally, there is previous work on diversity in network ranking. In this line of research, the goal is to find important nodes that also maximally cover the nodes in the network [31, 32].. Our problem is fundamentally different, since we look for rankings that follow a parity constraint.

BIBLIOGRAPHY

- [1] S. A. Friedler, C. Scheidegger, S. Venkatasubramanian, S. Choudhary, E. P. Hamilton, and D. Roth, “A comparative study of fairness-enhancing interventions in machine learning,” in *FAT**, 2019, pp. 329–338.
- [2] S. Brin and L. Page, “The anatomy of a large-scale hypertextual web search engine,” *Computer Networks and ISDN Systems*, vol. 30, no. 1, pp. 107 – 117, 1998.
- [3] J. M. Kleinberg, “Authoritative sources in a hyperlinked environment,” *J. ACM*, vol. 46, no. 5, p. 604–632, 1999.
- [4] R. Lempel and S. Moran, “Salsa: The stochastic approach for link-structure analysis,” *ACM Trans. Inf. Syst.*, vol. 19, no. 2, p. 131–160, 2001.
- [5] C. Castillo, D. Donato, A. Gionis, V. Murdock, and F. Silvestri, “Know your neighbors: Web spam detection using the web topology,” in *SIGIR*, 2007.
- [6] D. F. Gleich, “Pagerank beyond the web,” *SIAM Review*, vol. 57, no. 3, pp. 321–363, 2015.
- [7] T. H. Haveliwala, “Topic-sensitive pagerank,” in *WWW*, 2002.
- [8] C. Avin, B. Keller, Z. Lotker, C. Mathieu, D. Peleg, and Y. A. Pignolet, “Homophily and the glass ceiling effect in social networks,” in *ITCS*, 2015, pp. 41–50.
- [9] A. Stoica, C. J. Riederer, and A. Chaintreau, “Algorithmic glass ceiling in social networks: The effects of social recommendations on network diversity,” in *WebConf*, 2018, pp. 923–932.
- [10] V. Amelkin and A. K. Singh, “Fighting opinion control in social networks via link recommendation,” in *Proceedings of the 25th ACM SIGKDD International*

- Conference on Knowledge Discovery and Data Mining*, ser. KDD '19. New York, NY, USA: Association for Computing Machinery, 2019, p. 677–685. [Online]. Available: <https://doi.org/10.1145/3292500.3330960>
- [11] F. Masrour, T. Wilson, H. Yan, P. Tan, and A. Esfahanian, “Bursting the filter bubble: Fairness-aware network link prediction,” in *The Thirty-Fourth AAAI Conference on Artificial Intelligence, AAAI 2020, The Thirty-Second Innovative Applications of Artificial Intelligence Conference, IAAI 2020, The Tenth AAAI Symposium on Educational Advances in Artificial Intelligence, EAAI 2020, New York, NY, USA, February 7-12, 2020*. AAAI Press, 2020, pp. 841–848. [Online]. Available: <https://aaai.org/ojs/index.php/AAAI/article/view/5429>
- [12] M. Feldman, S. A. Friedler, J. Moeller, C. Scheidegger, and S. Venkatasubramanian, “Certifying and removing disparate impact,” in *KDD*, 2015, pp. 259–268.
- [13] C. Dwork, M. Hardt, T. Pitassi, O. Reingold, and R. S. Zemel, “Fairness through awareness,” in *Innovations in Theoretical Computer Science*, 2012, pp. 214–226.
- [14] K. S. Miller, “On the inverse of the sum of matrices,” *Mathematics Magazine*, vol. 54, no. 2, pp. 67–72, 1981. [Online]. Available: <https://doi.org/10.1080/0025570X.1981.11976898>
- [15] C. M. Grinstead and J. L. Snell, *Introduction to Probability*. AMS, 2003. [Online]. Available: http://www.dartmouth.edu/~chance/teaching_aids/books_articles/probability_book/book.html
- [16] R. A. Rossi and N. K. Ahmed, “The network data repository with interactive graph analytics and visualization,” in *AAAI*, 2015. [Online]. Available: <http://networkrepository.com>
- [17] L. A. Adamic and N. S. Glance, “The political blogosphere and the 2004 U.S. election: divided they blog,” in *LinkKDD*, 2005.
- [18] D. A. Easley and J. M. Kleinberg, *Networks, Crowds, and Markets*. Cambridge University Press, 2010.
- [19] Q. V. Liao and W. Fu, “Can you hear me now?: mitigating the echo chamber effect by source position indicators,” in *CSCW*, 2014, pp. 184–196.

- [20] A. Narayanan, “21 fairness definitions and their politics. (tutorial),” in *FAT*, 2018.
- [21] N. Mehrabi, F. Morstatter, N. Saxena, K. Lerman, and A. Galstyan, “A survey on bias and fairness in machine learning,” *CoRR*, vol. abs/1908.09635, 2019. [Online]. Available: <http://arxiv.org/abs/1908.09635>
- [22] M. Zehlike, F. Bonchi, C. Castillo, S. Hajian, M. Megahed, and R. Baeza-Yates, “Fa*ir: A fair top-k ranking algorithm,” in *CIKM*, 2017.
- [23] K. Yang and J. Stoyanovich, “Measuring fairness in ranked outputs,” in *SSDBM*, 2017.
- [24] A. J. Biega, K. P. Gummadi, and G. Weikum, “Equity of attention: Amortizing individual fairness in rankings,” in *SIGIR*, 2018, pp. 405–414.
- [25] A. Beutel, J. Chen, T. Doshi, H. Qian, L. Wei, Y. Wu, L. Heldt, Z. Zhao, L. Hong, E. H. Chi, and C. Goodrow, “Fairness in recommendation ranking through pairwise comparisons,” in *KDD*, 2019, pp. 2212–2220.
- [26] A. Singh and T. Joachims, “Fairness of exposure in rankings,” in *KDD*, 2018.
- [27] A. J. Bose and W. Hamilton, “Compositional fairness constraints for graph embeddings,” in *ICML*, 2019, pp. 715–724.
- [28] F. Masrour, T. Wilson, H. Yan, P.-N. Tan, and A.-H. Esfahanian, “Bursting the filter bubble: Fairness-aware network link prediction,” in *AAAI*, 2020.
- [29] T. A. Rahman, B. Surma, M. Backes, and Y. Zhang, “Fairwalk: Towards fair graph embedding,” in *IJCAI*, 2019, pp. 3289–3295.
- [30] F. Karimi, M. Génois, C. Wagner, P. Singer, and M. Strohmaier, “Homophily influences ranking of minorities in social networks,” *Nature Scientific Reports*, vol. 8, 2018.
- [31] X. Zhu, A. B. Goldberg, J. V. Gael, and D. Andrzejewski, “Improving diversity in ranking using absorbing random walks,” in *HLT-NAACL*, 2007, pp. 97–104.
- [32] Q. Mei, J. Guo, and D. R. Radev, “Divrank: the interplay of prestige and diversity in information networks,” in *KDD*, 2010, pp. 1009–1018.

- [33] L. Takac and M. Zábovský, “Data analysis in public social networks,” in *International Scientific Conference and International Workshop Present Day Trends of Innovations, Lomza, Poland*, 2012.
- [34] J. Tang, J. Zhang, L. Yao, J. Li, L. Zhang, and Z. Su, “Arnetminer: extraction and mining of academic social networks,” in *KDD*, 2008, pp. 990–998.
- [35] Y. Zhang, J. Tang, Z. Yang, J. Pei, and P. S. Yu, “COSNET: connecting heterogeneous social networks with local and global consistency,” in *KDD*, 2015, pp. 1485–1494.

APPENDIX A

PROOFS

A.1 Fairness Aware PageRank Ranking

A.2 PageRankFair Recommendations

A.1 Fairness Aware PageRank Ranking

When $\phi = r$, the average weight of red nodes is equal with the average weight of the blue nodes, i.e., $\frac{\sum_{v \in R} w_v}{|R|} = \frac{\sum_{v \in B} w_v}{|B|}$.

Proof. It holds: $\frac{\sum_{v \in B} w_v}{|B|} = \frac{\sum_{v \in V} w_v - \sum_{v \in R} w_v}{|N| - |R|} =$
 $\frac{1/r \sum_{v \in R} w_v - \sum_{v \in R} w_v}{|N| - |R|} =$
 $\frac{N/|R| \sum_{v \in R} w_v - \sum_{v \in R} w_v}{|N| - |R|} = \frac{\sum_{v \in R} w_v}{|R|}.$ □

Proof of Lemma 3.2.

Proof. From the transition matrix \mathbf{P}_L , each node $i \in L_R$ gives a portion $\frac{1-\phi}{out_B(i)}$ of each of its pagerank to its neighbors. The blue neighbors do not get any residual pagerank, thus they get an $1 - \phi$ portion as in the LFPR_N algorithm. Each of the red neighbors gets an additional $\frac{1}{out_R(i)} \delta_R(i) = \frac{1}{out_R(i)} (\phi - \frac{(1-\phi) out_R(i)}{out_B(i)})$, which sums to $\frac{\phi}{out_R(i)}$. Thus, the red nodes get an ϕ portion of i 's pagerank as in the LFPR_N algorithm. The proof is analogous for each node $i \in L_B$. □

A lower bound for the optimal weight redistribution for targeted fairness

Given the weight vector \mathbf{w} and the set S , we divide the full set of nodes in three categories: the set B_S of blue nodes in S , the set R_S of red nodes in S , and the rest of nodes O not in S . In order for \mathbf{f} to be fair, it must be that it moves weight between these categories. Furthermore, this movement is always in one direction, e.g., all nodes in R_S will increase their weight. It is clearly sub-optimal to increase the weight of some nodes in R_S and decrease the weight of others. We define the variables $x_B = \sum_{i \in B_S} (f_i - w_i)$, $x_R = \sum_{i \in R_S} (f_i - w_i)$, and $x_O = \sum_{i \in O} (f_i - w_i)$ to be the total change in weight for the nodes in B_S , R_S , and O respectively. Note that these values may be positive, indicating an increase in weight for the respective category, or negative, indicating a decrease in weight for the respective category. It holds:

$$x_B + x_R + x_O = 0 \quad (\text{A.1})$$

Let \mathbf{f}_R and \mathbf{f}_B the weight allocated to nodes in R_S and B_S respectively, and ρ and β be their desired values according to ϕ . Also, let w_B and w_R be the weight of the nodes in B_S and R_S respectively. Since the vector \mathbf{f} is fair for the nodes in S it holds that

$$\frac{w_R + x_R}{w_B + x_B} = \frac{\rho}{\beta} \quad (\text{A.2})$$

Using Equations A.1 and A.2, we can express x_B and x_R as a function of x_O :

$$x_R = \rho w_B - \beta w_R - \rho x_O \quad (\text{A.3})$$

$$x_B = \beta w_R - \rho w_B - \beta x_O \quad (\text{A.4})$$

Now, let N_B , N_R , and N_O denote the number of nodes in categories B_S , R_S , and O respectively. To minimize loss, and since we allow \mathbf{f} to have negative entries, the change in weight must be distributed equally in each category. Thus, the total loss is

$$\text{Loss}(\mathbf{f}, \mathbf{w}) = \frac{x_R^2}{N_R} + \frac{x_B^2}{N_B} + \frac{x_O^2}{N_O} \quad (\text{A.5})$$

A.2 PageRankFair Recommendations

APPENDIX B

EXPERIMENT EVALUATION

B.1 Fairness Aware PageRank Ranking

B.2 PageRankFair Recommendations

B.1 Fairness Aware PageRank Ranking

Convert figures to compatible form

B.1.1 Reproducibility

We substitute Equations A.3 and A.4 in Equation A.5, we take the derivative with respect to x_O , and we set it zero. Solving for x_O , we get:

Table B.1: Real dataset characteristics. r , b relative size of protected and unprotected group, respectively; p_R , p_B pagerank assigned to the red and blue group respectively

Dataset	#nodes	#edges	Protected attribute	r	b	homophily	p_R	p_B
POKEC	1,632,803	30,622,564	gender (women)	0.51	0.49	1.11	0.54	0.46
DBLP1	423,469	2,462,422	gender (women)	0.19	0.81	0.83	0.13	0.87
LINKEDIN	3,209,448	13,016,453	gender (women)	0.37	0.63	0.72	0.37	0.63
PHYSICS	30,359	347,235	year (after 1997)	0.66	0.34	0.76	0.39	0.61

Table B.2: Utility loss with respect to optimal utility ($\frac{LFPR_X}{OPTIMAL}$)

Dataset	LFPR_N	LFPR_U	LFPR_P	SFPR
POKEC	30.57	35.31	15.39	-
TWITTER	152.08	156.02	66.92	6.94
DBLP1	99.12	41.66	21.80	-
DBLP2	95.84	47.81	25.18	6.13
LINKEDIN	4,913	1,787	1,149	-
PHYSICS	9.56	9.04	8.21	50.24

$$x_O = \frac{N_O(\beta w_B - \rho w_R)(\beta N_R - \rho N_B)}{\rho N_B(\rho N_O + N_R) + \beta N_R(\beta N_O + N_B)} \quad (\text{B.1})$$

Substituting x_O in Equations A.4 and A.3, we obtain:

$$x_R = \frac{(\rho w_B - \beta w_R)N_R(\beta N_O + N_B)}{\rho N_B(\rho N_O + N_R) + \beta N_R(\beta N_O + N_B)} \quad (\text{B.2})$$

$$x_B = \frac{(\beta w_R - \rho w_B)N_B(\rho N_O + N_R)}{\rho N_B(\rho N_O + N_R) + \beta N_R(\beta N_O + N_B)} \quad (\text{B.3})$$

There are some interesting observations in these equations. First, a factor that appears in all equations is $\beta \mathbf{w}_R - \rho \mathbf{w}_B$, which tells us how unfair the original weights are. For example, if $\beta \mathbf{w}_R - \rho \mathbf{w}_B < 0$, then we are unfair towards category R . In this case the nodes in category R will always receive weight $w_R > 0$. The origin of the weight depends on the ratio N_R/N_B of the nodes in S . If $\beta N_R - \rho N_B < 0$, then we have proportionally more nodes of B in S with an excess of weight. In this case we remove weight only from the nodes in B , and we distribute it to the nodes in R and O as defined by Equations B.3 and B.1. If $\beta N_R - \rho N_B > 0$, then we have proportionally less nodes of B in S , but they have proportionally more weight. In this case we remove weight from both the nodes in B , and O , as defined by Equations B.3 and B.1, and we distribute it to the nodes in R . If $\beta N_R - \rho N_B = 0$, then we take weight only from the nodes in B and give only to the nodes in R .

Having computed the values for x_R , x_B and x_O , we can now compute the loss using Equation A.5. Note that this is a lower bound to the optimal loss for our problem, since it does not guarantee that the resulting vector \mathbf{f} has non-negative entries.

Additional datasets and experiments

In Table B.1, we present statistics for additional datasets.




Figure B.1: Cutting Point for Selecting Nodes.

- **POKEC** [33]: This is a Slovak social network. Nodes correspond to users, and links to friendships. Friendship relations are directed.
- **DBLP1**: An author collaboration network constructed by the Arnetminer academic search system [34] using publication data from dblp. Two authors are connected if they have co-authored an article.
- **LINKEDIN** [35]: Nodes correspond to LinkedIn profiles. Two profiles are linked if they were co-viewed by the same user.
- **PHYSICS**: This is the Arxiv HEP-PH (high energy physics phenomenology) citation graph from the SNAP dataset¹. Nodes correspond to papers and there is an edge from a paper to another, if the first paper cites the second one.

Again, there are cases where the fraction of the weight assigned to the protected group is even smaller than r .

In Figure ??, we report results for the original and the locally fair PageRank algorithms for the additional datasets and in Figure ??, we report results for the locally fair PageRank algorithms for $\phi = r$.

B.2 PageRankFair Recommendations

Similar results with those in analysis for random source nodes can be derived for the other two source nodes sets from the tables B.3, B.4, B.5, B.6. Selection threshold for each set is presented in figures B.1, B.2, B.3, B.4.

¹<http://snap.stanford.edu/data>

Table B.3: Target Quality Features in Blogs - Red Source Nodes.

Policy	Distance			PageRank			Red PageRank			Node Homophily		
	mean	median	max	mean	median	max	mean	median	max	mean	median	max
Random	3.380000	3	5	0.000822	0.000339	0.045172	0.332817	0.282878	0.638946	0.435733	0.500000	1.000000
Node2vec	2.722222	3	4	0.004934	0.004793	0.010006	0.340760	0.321328	0.564971	0.308273	0.160000	0.957143
Gain	3.890000	4	7	0.000284	0.000243	0.000583	0.622608	0.620985	0.638946	1.000000	1.000000	1.000000
ExpGain	3.040816	3	5	0.000940	0.000583	0.002620	0.580573	0.590254	0.638946	0.969143	1.000000	1.000000

Table B.4: Target Quality Features in Blogs - Blue Source Nodes.


Policy	Distance			PageRank			Red PageRank			Node Homophily		
	mean	median	max	mean	median	max	mean	median	max	mean	median	max
Random	2.846667	3	6	0.000840	0.000331	0.045172	0.337866	0.282878	0.638946	0.445203	0.500000	1.000000
Node2vec	2.437037	2	4	0.004934	0.004793	0.010006	0.340760	0.321328	0.564971	0.308273	0.160000	0.957143
Gain	4.200000	4	7	0.000284	0.000243	0.000583	0.622608	0.620985	0.638946	1.000000	1.000000	1.000000
ExpGain	3.160839	3	5	0.000982	0.000742	0.002620	0.578018	0.577612	0.638946	0.967215	0.989131	1.000000

Table B.5: Target Quality Features in Twitter - Red Source Nodes.


Policy	Distance			PageRank			Red PageRank			Node Homophily		
	mean	median	max	mean	median	max	mean	median	max	mean	median	max
Random	4.762500	5	7	0.000054	0.000037	0.001418	0.579776	0.639552	1.000000	0.511345	0.500000	1.000000
Node2vec	3.636364	4	7	0.001447	0.001376	0.003275	0.590534	0.721374	0.817765	0.650000	1.000000	1.000000
Gain	4.775000	5	7	0.000185	0.000228	0.000298	0.942876	1.000000	1.000000	1.000000	1.000000	1.000000
ExpGain	4.442857	4	7	0.000457	0.000283	0.001412	0.935682	1.000000	1.000000	1.000000	1.000000	1.000000

Table B.6: Target Quality Features in Twitter - Blue Source Nodes.


Policy	Distance			PageRank			Red PageRank			Node Homophily		
	mean	median	max	mean	median	max	mean	median	max	mean	median	max
Random	4.754545	5	7	0.000058	0.000037	0.001418	0.575291	0.639552	1.000000	0.511240	0.500000	1.000000
Node2vec	3.829545	4	5	0.001574	0.001376	0.003275	0.651073	0.736853	0.817765	0.750000	1.000000	1.000000
Gain	5.372727	5	7	0.000185	0.000228	0.000298	0.942876	1.000000	1.000000	1.000000	1.000000	1.000000
ExpGain	4.904255	5	7	0.000457	0.000283	0.001412	0.935682	1.000000	1.000000	1.000000	1.000000	1.000000



(a) Node2vec




(b) Fair




(c) Hybrid Fair


Figure B.2: Cutting Point for Selecting Nodes.



(a) Node2vec




(b) Fair




(c) Hybrid Fair


Figure B.3: Cutting Point for Selecting Nodes.



(a) Node2vec



(b) Fair



(c) Hybrid Fair

Figure B.4: Cutting Point for Selecting Nodes.

AUTHOR'S PUBLICATIONS

Προαιρετικά, βάζουμε μία λίστα με τις δημοσιεύσεις του συγγραφέα.

SHORT BIOGRAPHY

Ένα σύντομο βιογραφικό είναι απαραίτητο.